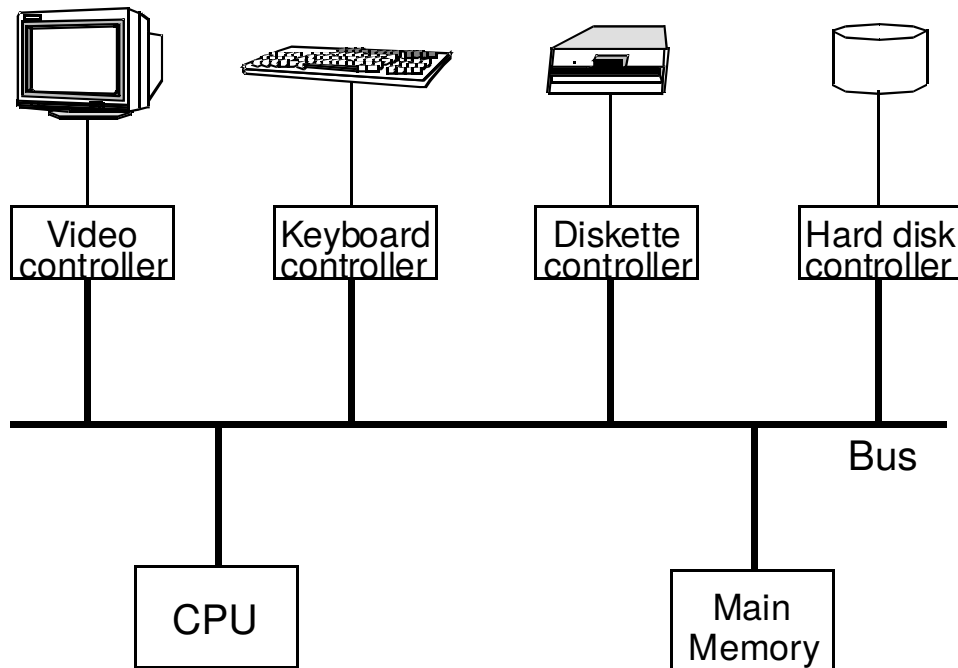


TDDI11: Embedded Software

Input/Output Programming


I/O device controller



- Controller
 - controls its I/O device and handles bus access for it.
- Example
 - Hard disk controller receives, for example, a read request from the CPU \Rightarrow gives corresponding commands to the device in order to execute the request, collects data and organises the incoming bit stream into units to be sent on the bus.
- A controller has to interact with the bus or CPU, according to a certain protocol, in order to transmit and/or receive.

Microcontroller


- Processor, memory, a clock oscillator, I/O, ADC, DAC, possibly on the same chip.



How is the I/O connected to microcontroller?


Parallel I/O

- Processor has one or more N-bit ports
- Processor's software reads and writes a port just like a register
 - E.g., PA = 0xF2;





Bus-based I/O

- Processor has ports that form a single bus
- Communication protocol is built into the processor
- A single instruction for reading or writing



Example: Serial Transmission Using UARTs

- UART: Universal Asynchronous Receiver Transmitter
 - Takes parallel data and transmits serially
 - Receives serial data and converts to parallel
- Parity: extra bit for simple error checking
- Start bit, stop bit
- Baud rate
 - signal changes per second
 - bit rate usually higher




Serial Transmission Using UARTs, cont.

- Transmitter
 - Host process writes in the transmit buffer of UART
 - Sends “start bit” over **transmit pin (tx)**
 - Signals beginning of transmission to the remote UART
 - Shifts out the data in the buffer over **tx** on a specified rate
 - Transmits (optional) parity bit
 - Signals host processor indicating that is ready to transmit more
- Receiver
 - Constantly monitoring the **receive pin (rx)** for a “start bit”
 - Receiver starts sampling the **rx** pin at predetermined intervals
 - Shifts sampled bits into the receive shift register
 - Reads (optional) parity bit
 - Once data is received, it signals host processor
 - Host processor reads the byte out of the receive shift register
 - The receiver is ready for more data

Parallel I/O peripheral


- When processor only supports bus-based I/O but parallel I/O needed
- Each port on peripheral connected to a register within peripheral that is read/written by the processor



Adding parallel I/O to a bus-based I/O processor

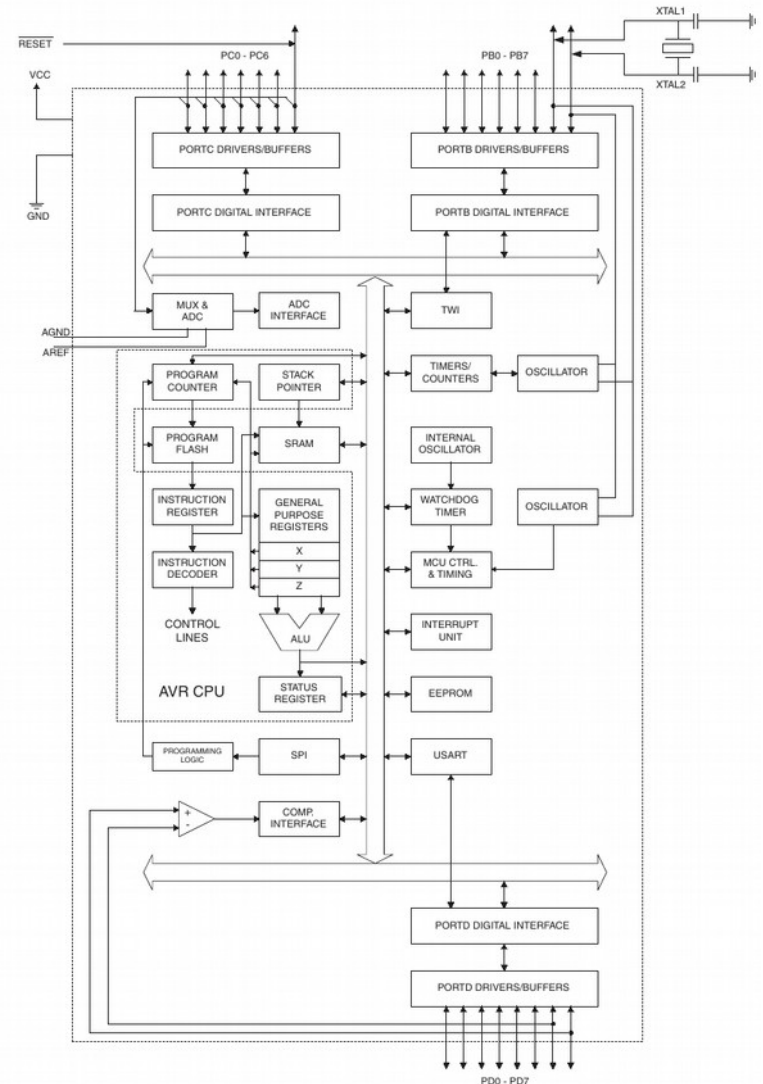
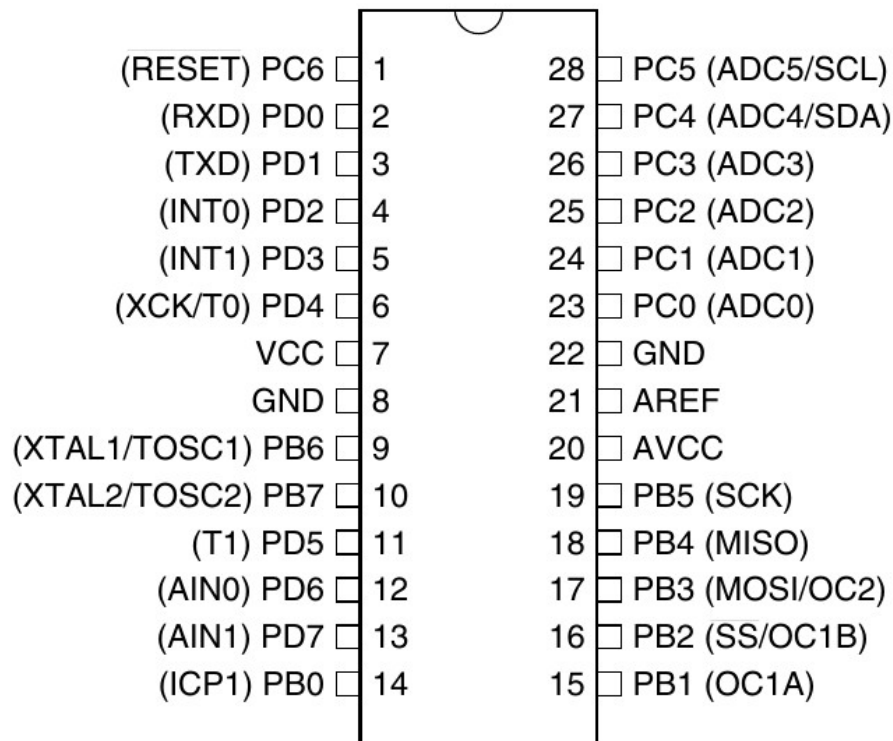
Extended parallel I/O

- When processor supports port-based I/O but more ports needed
- One or more processor ports interface with parallel I/O peripheral extending total number of ports available for I/O
 - e.g., extending 4 ports to 6 ports in figure



Extended parallel I/O

Example: ATmega8



How does a programmer talk to I/O?

Talk to peripherals: memory-mapped I/O and standard I/O

- Memory-mapped I/O
 - Peripheral registers occupy addresses in same address space as memory
 - e.g., Bus has 16-bit address
 - lower 32K addresses may correspond to memory
 - upper 32k addresses may correspond to peripherals
- Standard I/O (also called Isolated or I/O-mapped I/O)
 - Separate address lines for I/O and for memory
 - (Or) Additional pin (*M/IO*) on bus indicates whether a memory or peripheral access
 - e.g., Bus has 16-bit address
 - all 64K addresses correspond to memory when *M/IO* set to 0
 - all 64K addresses correspond to peripherals when *M/IO* set to 1

Memory-mapped I/O vs. Standard I/O

Memory-mapped I/O

- No special instructions
- Assembly instructions involving memory like MOV and ADD work with peripherals as well

Standard I/O

- No loss of memory addresses to peripherals
- Simpler address decoding logic in peripherals
- Small number of devices
-> high-order address bits can be ignored (smaller and/or faster comparators)

Principles to interface with I/O

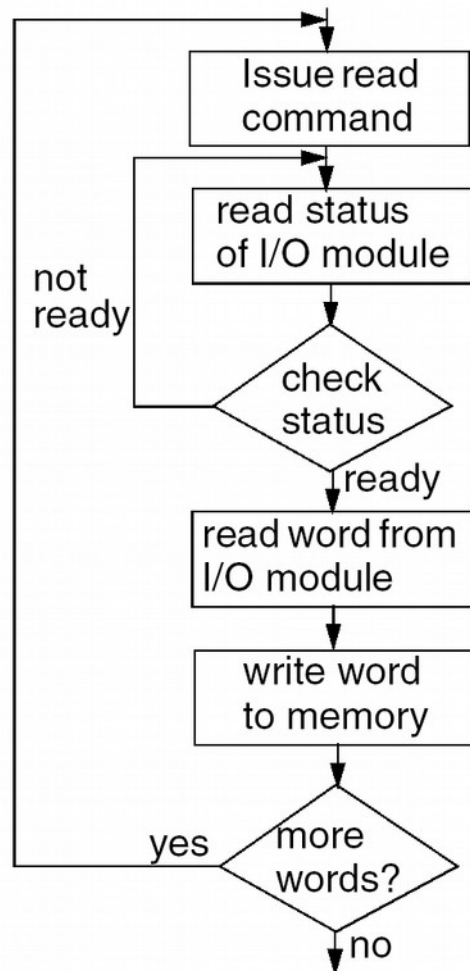
I/O Processing

- Three techniques are possible for transferring data to/from I/O devices:
 1. Programmed I/O
 2. Interrupt-driven I/O
 3. Direct memory access

I/O Processing

- Three techniques are possible for transferring data to/from I/O devices:
 1. **Programmed I/O**
 2. Interrupt-driven I/O
 3. Direct memory access

Programmed I/O (Polling)

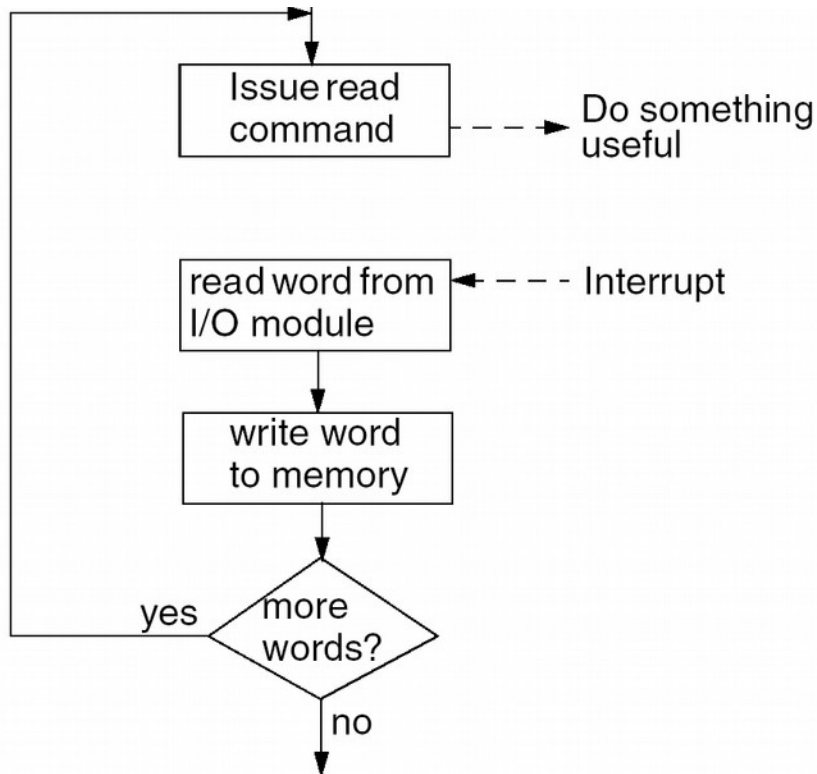


- The CPU executes a sequence of instructions, being in direct control of the I/O operations (sensing device status, read/write commands, etc.). When the CPU issues a command to the I/O module, it must wait until the I/O operation is complete.
- A lot of wasted time, because the CPU is much faster than devices.
- Sometimes preferred in real-time embedded systems!

I/O Processing

1. Programmed I/O
- 2. Interrupt-driven I/O**
3. Direct memory access

Interrupt-driven I/O



- After issuing an I/O command, the CPU need not wait until the operation has finished; instead of waiting, the CPU continues with other useful work.
- When the I/O operation has been completed, the I/O module issues an interrupt signal on the bus.
- After receiving the interrupt, the CPU moves the data to/from memory, and issues a new command if more data has to be read/written.

Interrupts

- Interrupt
 - Asynchronous electrical signal from a peripheral to the processor
- Consists of:
 - Interrupt pin (on the outside of the processor chip)
 - If Int is 1, processor suspends current program, jumps to an Interrupt Service Routine, or ISR
 - Interrupt service routine (ISR)
 - Small piece of software executed by the processor when an interrupt is asserted
 - Programmer must write and “install” the ISR

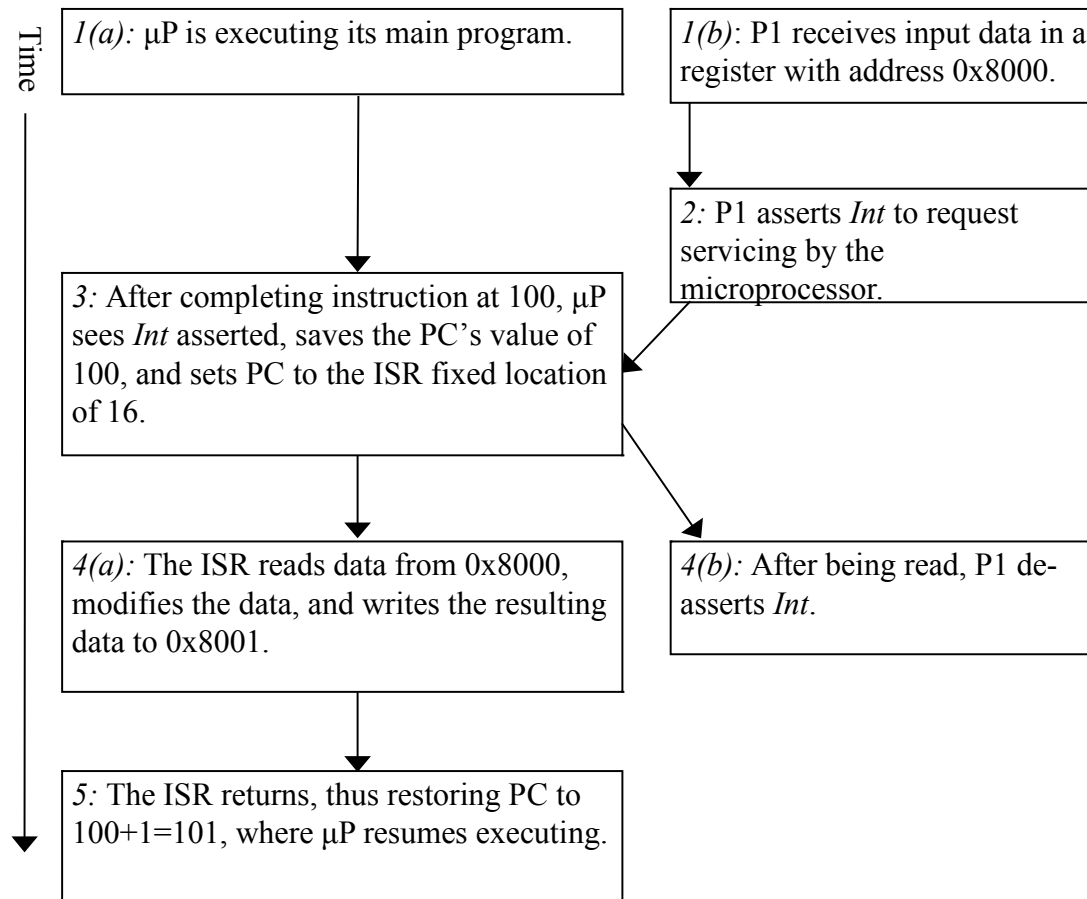
Interrupts, cont.

- What is the address (interrupt address vector) of the ISR?
 - Fixed interrupt
 - Address built into microprocessor, cannot be changed
 - Either ISR stored at address or a jump to actual ISR stored if not enough bytes available
 - Vectored interrupt
 - Peripheral must provide the address
 - Common when microprocessor has multiple peripherals connected by a system bus
 - Compromise: interrupt address table

Interrupts, cont.

- What is the address of the ISR?
 - **Fixed interrupt**
 - Vectored interrupt
 - Compromise: interrupt address table

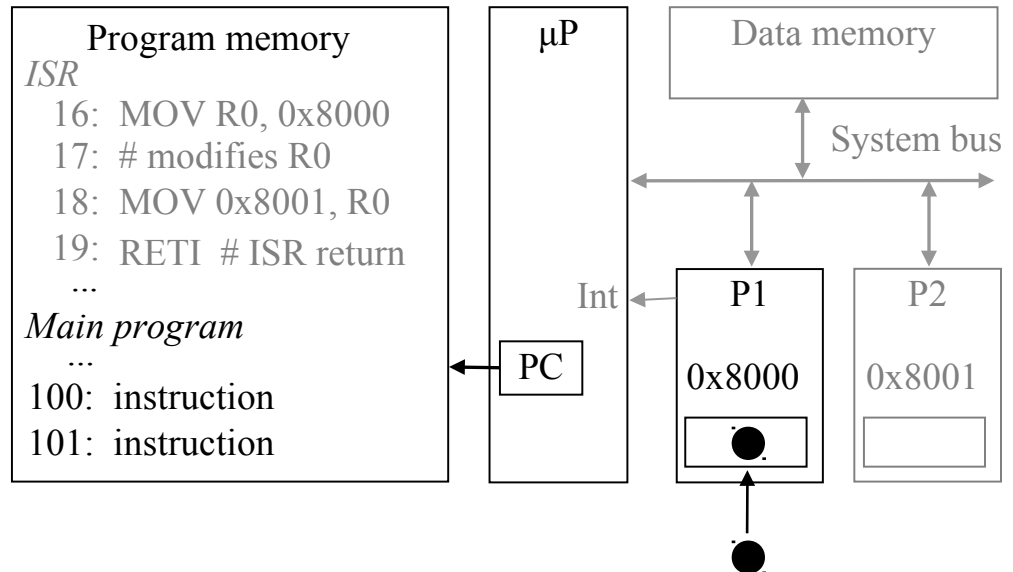
Interrupt-driven I/O using fixed ISR location



Interrupt-driven I/O using fixed ISR location

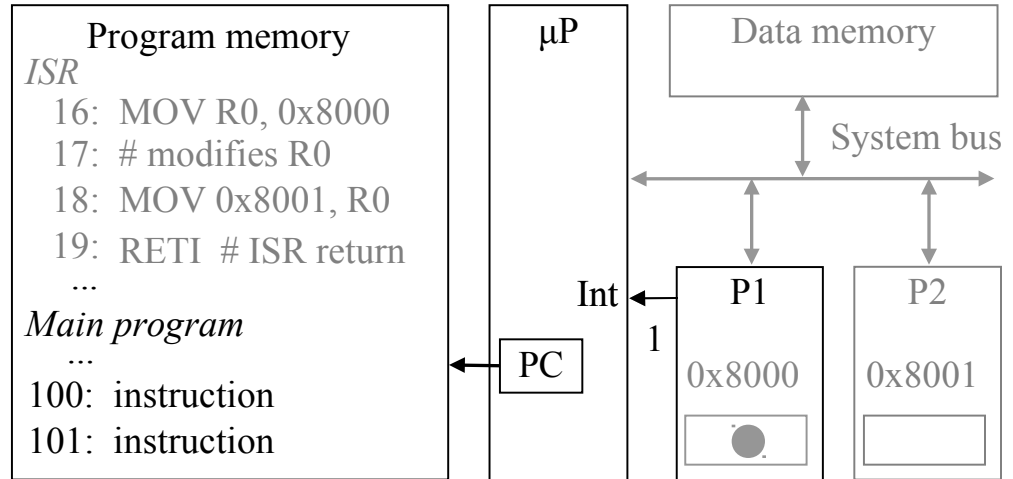
1(a): μP is executing its main program

1(b): P1 receives input data in a register with address 0x8000.



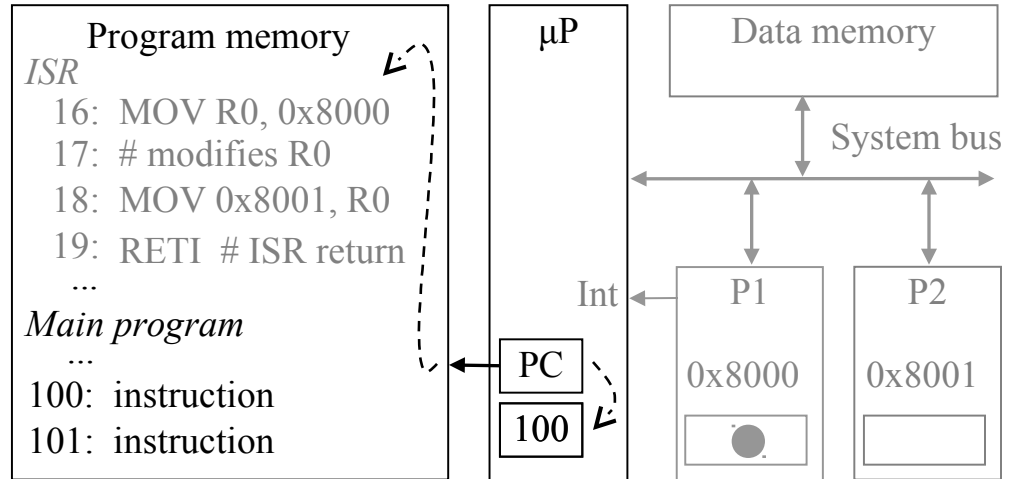
Interrupt-driven I/O using fixed ISR location

2: P1 asserts *Int* to request servicing by the microprocessor



Interrupt-driven I/O using fixed ISR location

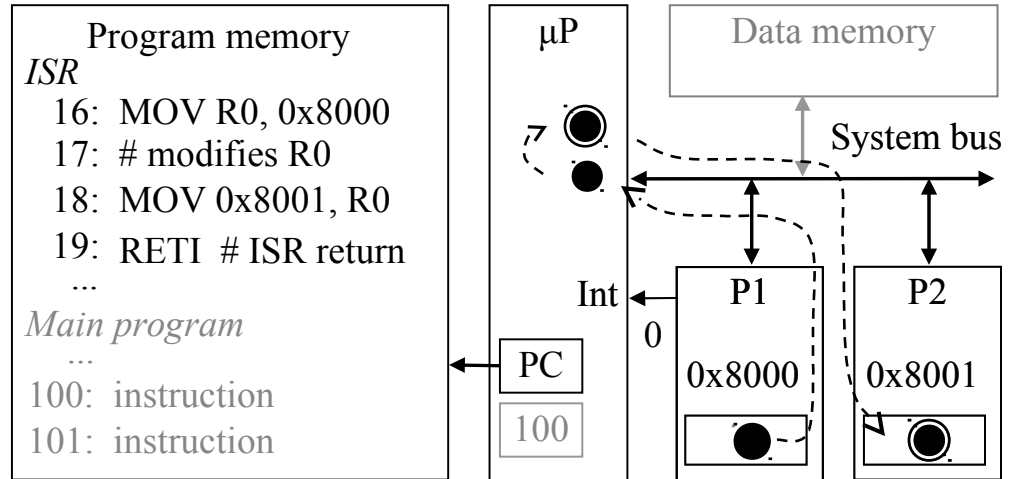
3: After completing instruction at 100, μP sees *Int* asserted, saves the PC's value of 100, and sets PC to the ISR fixed location of 16.



Interrupt-driven I/O using fixed ISR location

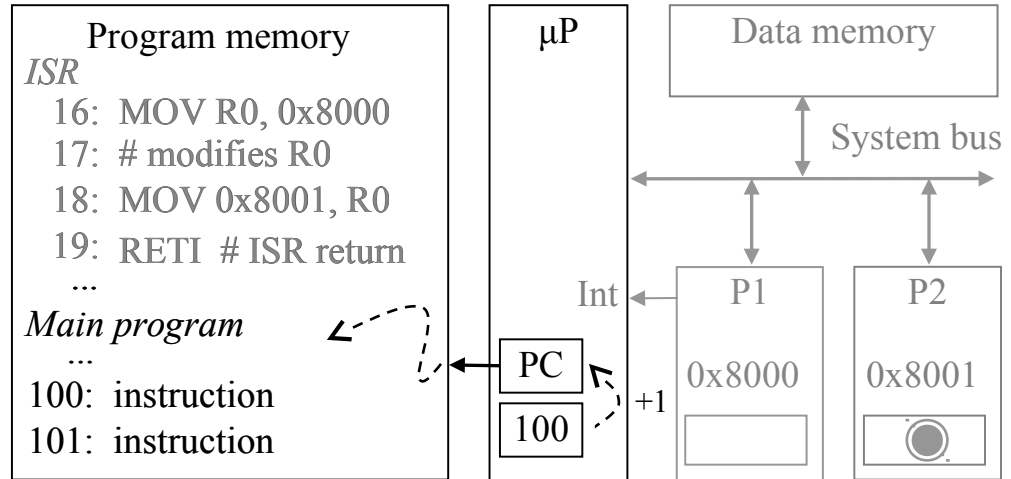
4(a): The ISR reads data from 0x8000, modifies the data, and writes the resulting data to 0x8001.

4(b): After being read, P1 deasserts *Int*.



Interrupt-driven I/O using fixed ISR location

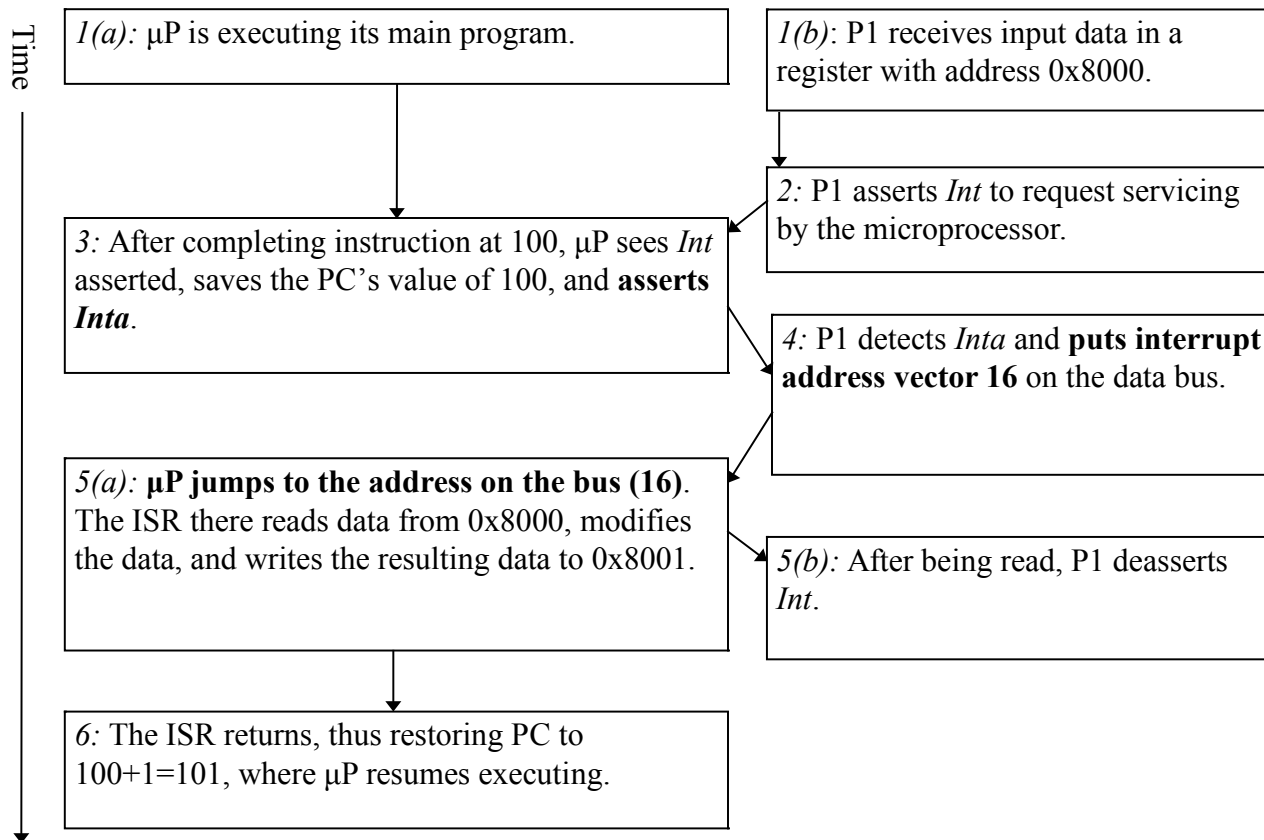
5: The ISR returns, thus restoring PC to $100+1=101$, where μP resumes executing.



Interrupts, cont.

- What is the address of the ISR?
 - Fixed interrupt
 - **Vectored interrupt**
 - Compromise: interrupt address table

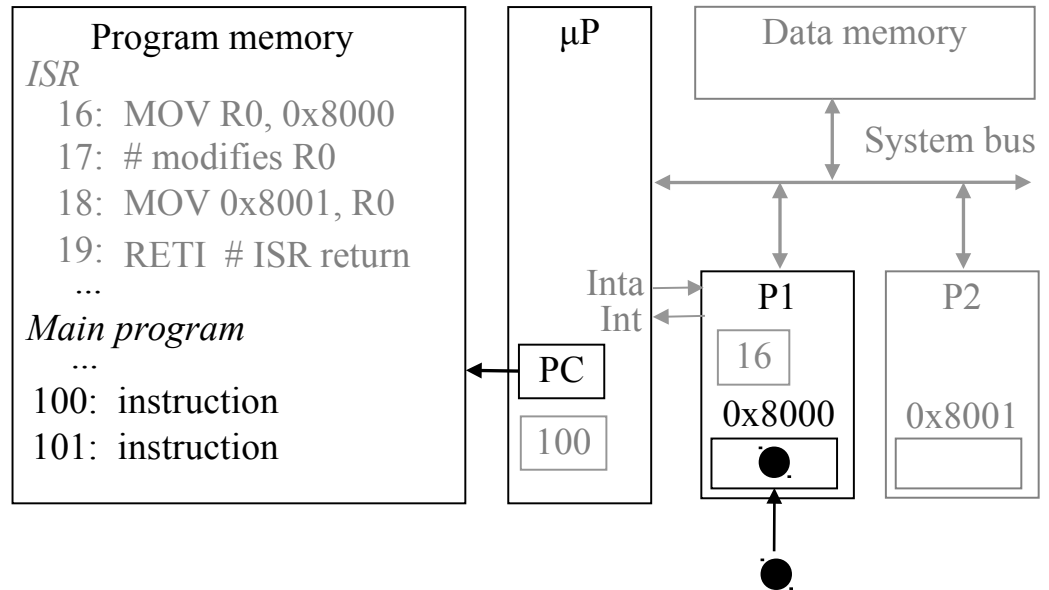
Interrupt-driven I/O using vectored interrupt



Interrupt-driven I/O using vectored interrupt

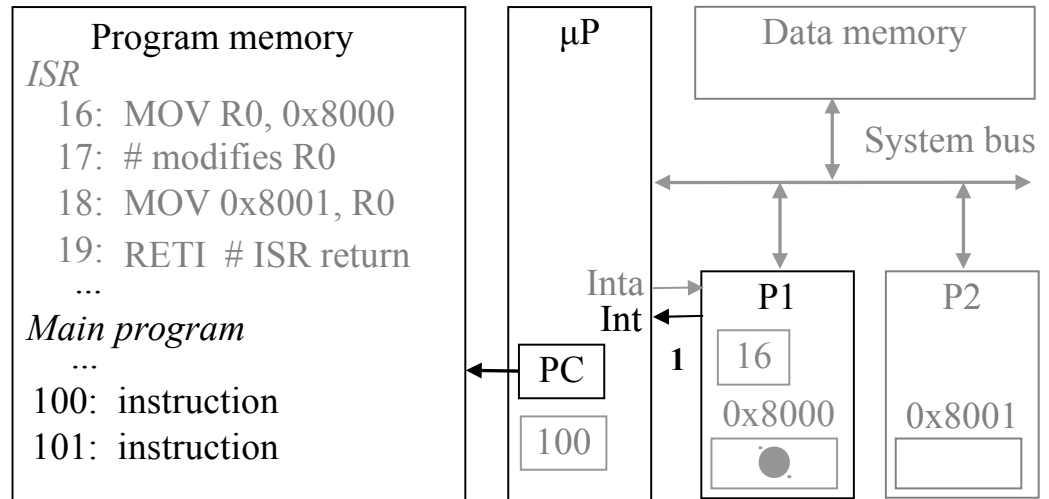
1(a): P is executing its main program

1(b): P1 receives input data in a register with address 0x8000.



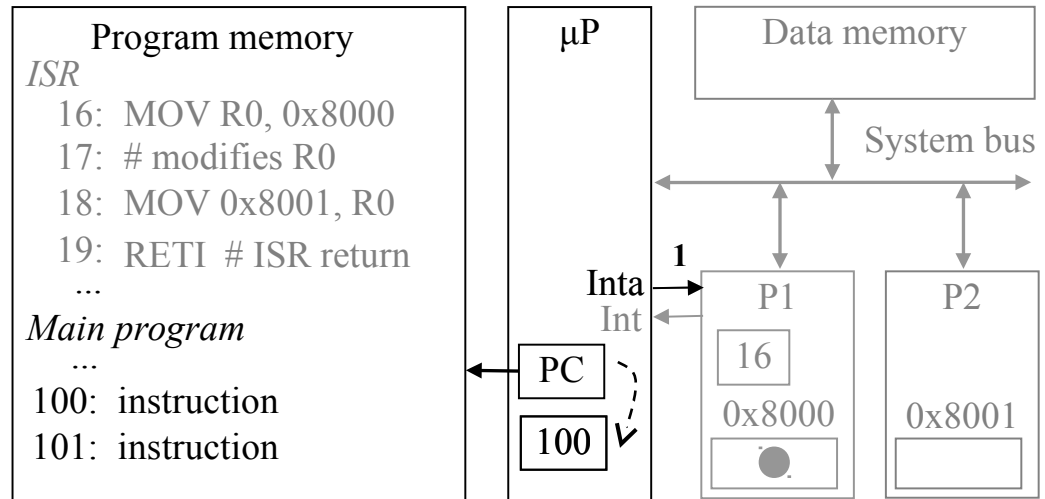
Interrupt-driven I/O using vectored interrupt

2: P1 asserts *Int* to request servicing by the microprocessor



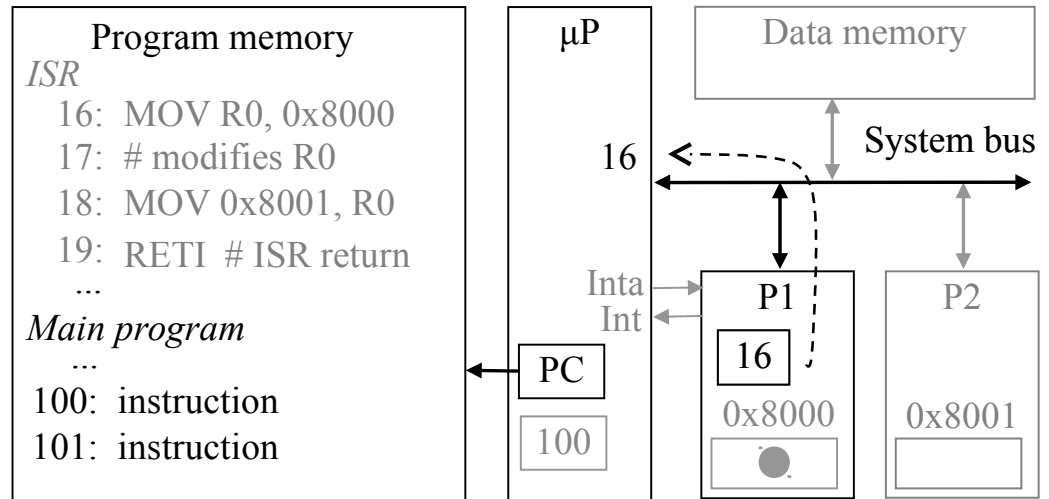
Interrupt-driven I/O using vectored interrupt

3: After completing instruction at 100, μP sees *Int* asserted, saves the PC's value of 100, and asserts *Inta*



Interrupt-driven I/O using vectored interrupt

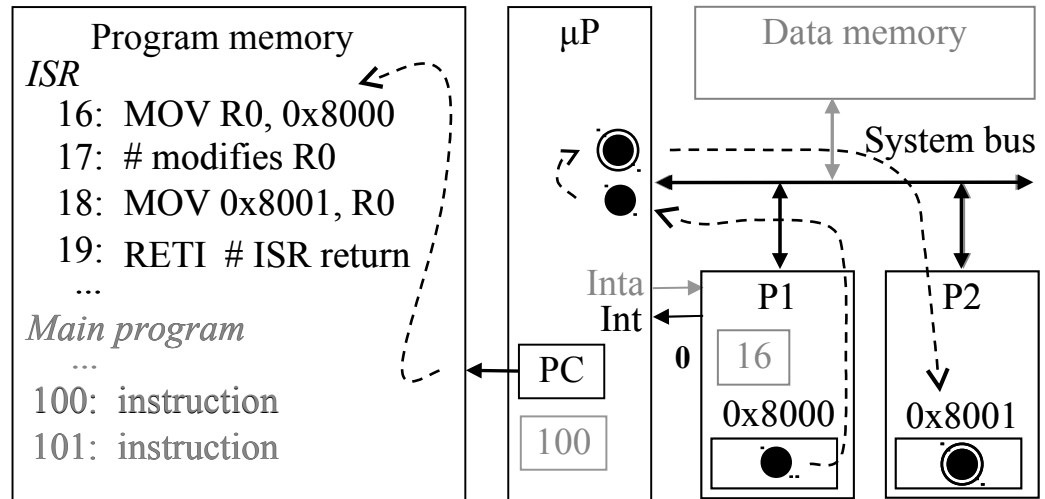
4: P1 detects *Inta* and puts **interrupt address vector 16** on the data bus



Interrupt-driven I/O using vectored interrupt

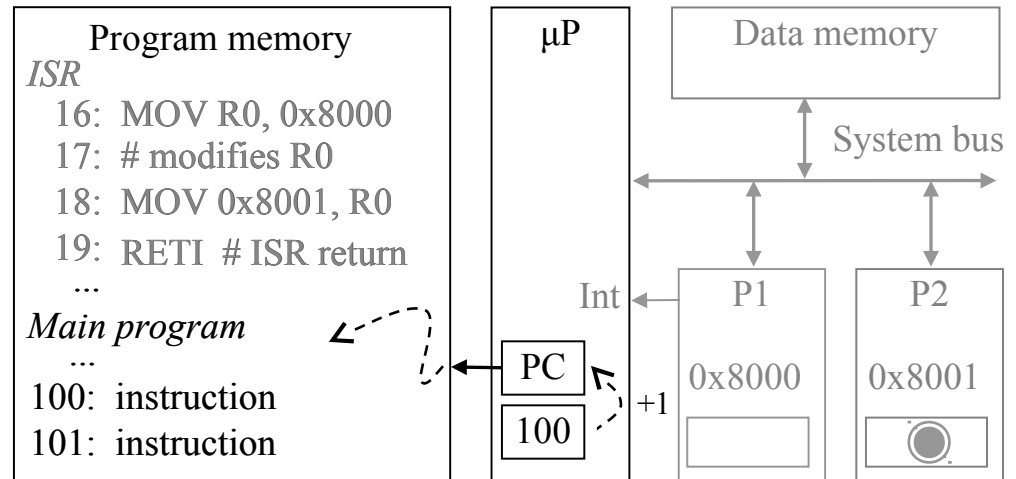
5(a): PC jumps to the address on the bus (16). The ISR there reads data from 0x8000, modifies the data, and writes the resulting data to 0x8001.

5(b): After being read, P1 deasserts *Int*.



Interrupt-driven I/O using vectored interrupt

6: The ISR returns, thus restoring the PC to $100+1=101$, where the μP resumes



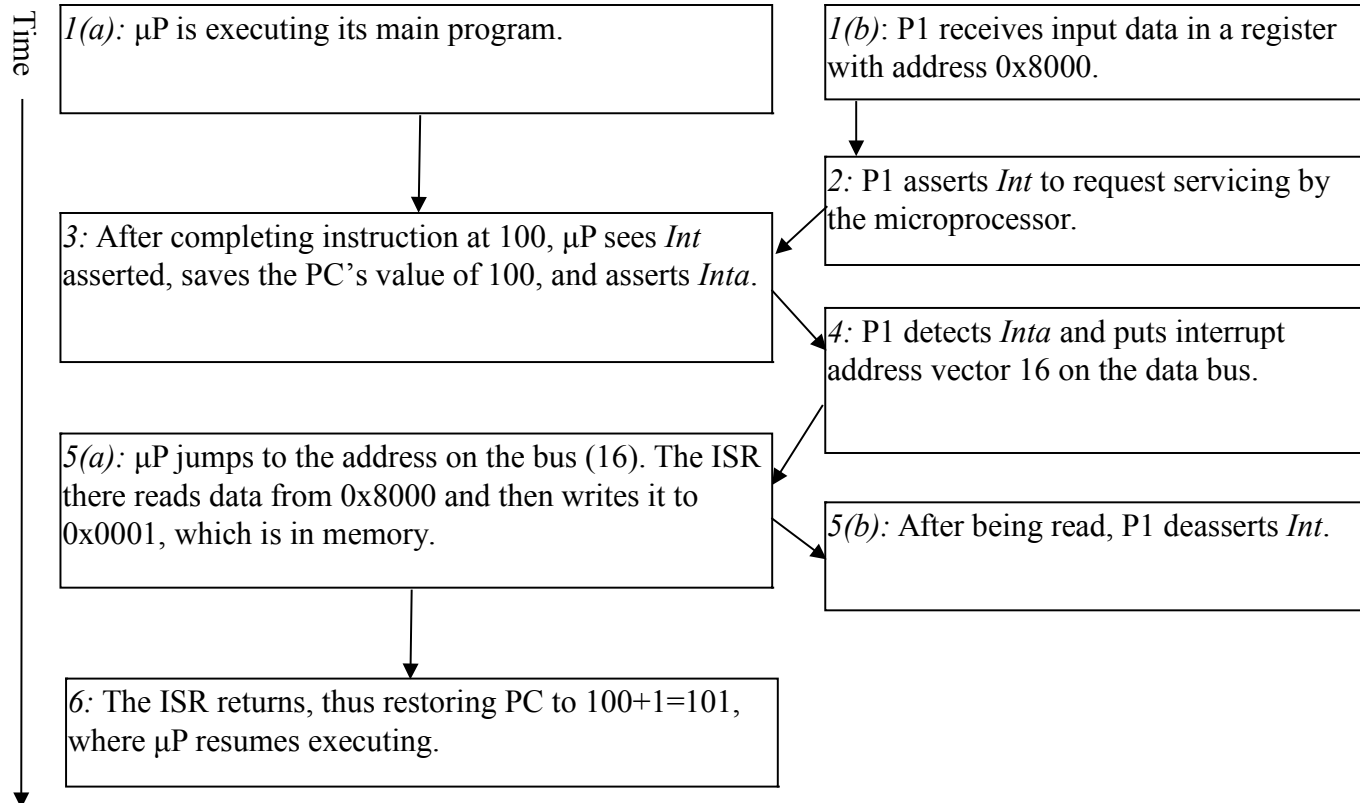
Interrupts, cont.

- What is the address of the ISR?
 - Fixed interrupt
 - Vectored interrupt
 - **Compromise: interrupt address table**

Interrupt address table

- Compromise between fixed and vectored interrupts
 - One interrupt pin
 - Table in memory holding ISR addresses (maybe 256 words)
 - Peripheral doesn't provide ISR address, but rather index into table
 - Fewer bits are sent by the peripheral
 - Can move ISR location without changing peripheral

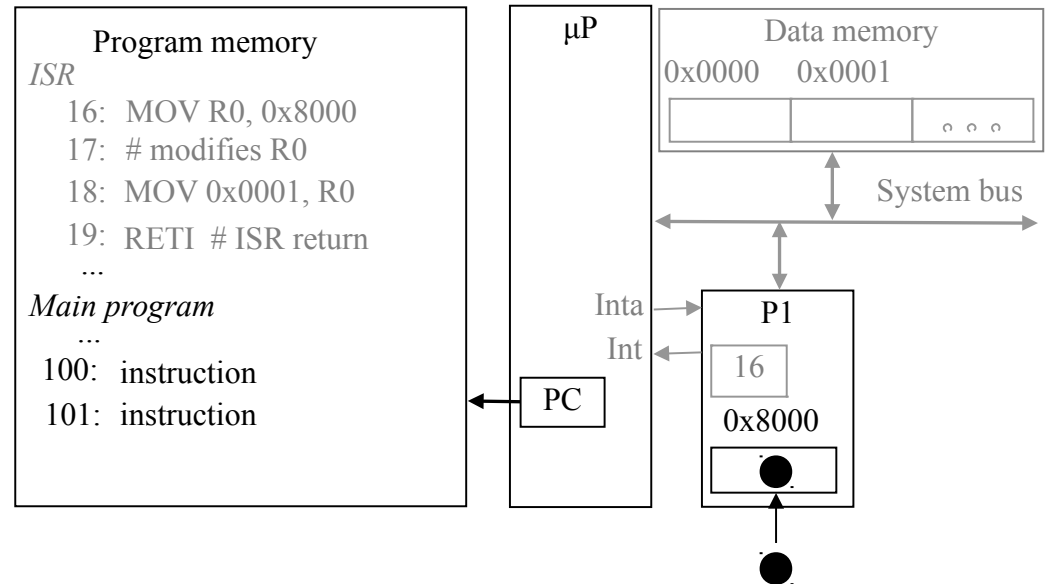
Peripheral to memory transfer *without* DMA, using vectored interrupt



Peripheral to memory transfer *without* DMA, using vectored interrupt

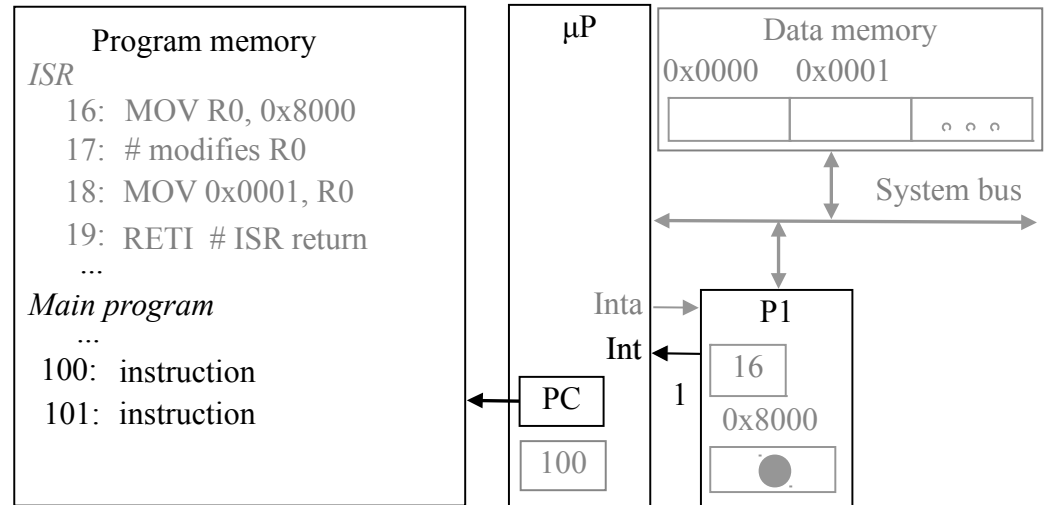
1(a): μP is executing its main program

1(b): P1 receives input data in a register with address 0x8000.



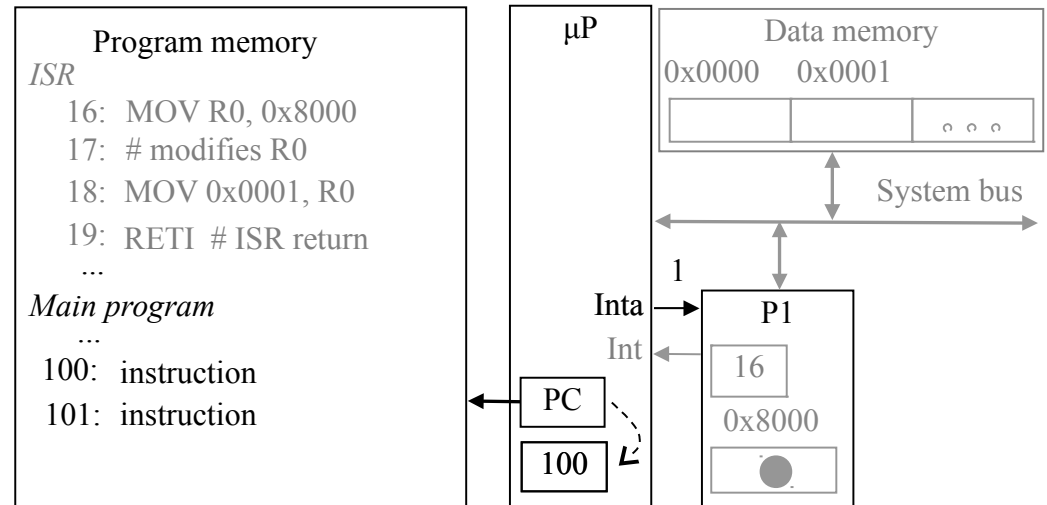
Peripheral to memory transfer *without* DMA, using vectored interrupt

2: P1 asserts *Int* to request servicing by the microprocessor



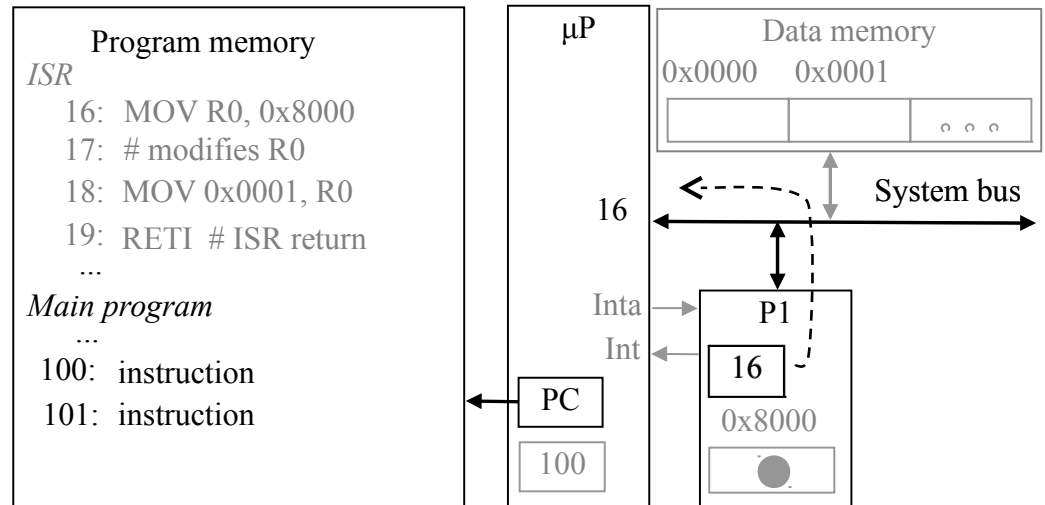
Peripheral to memory transfer *without* DMA, using vectored interrupt

3: After completing instruction at 100, μP sees *Int* asserted, saves the PC's value of 100, and asserts *Inta*.



Peripheral to memory transfer *without* DMA, using vectored interrupt (cont')

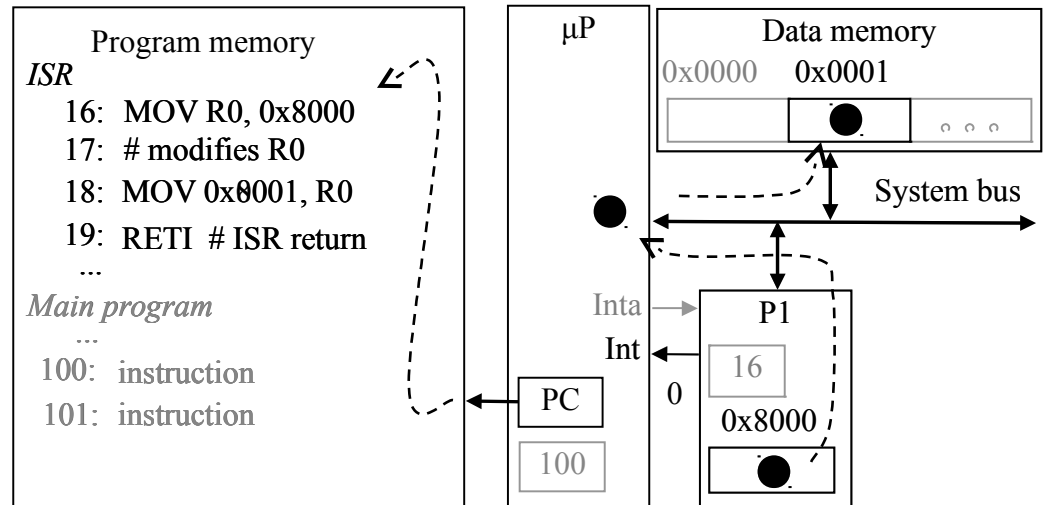
4: P1 detects *Inta* and puts interrupt address vector 16 on the data bus.



Peripheral to memory transfer *without* DMA, using vectored interrupt (cont')

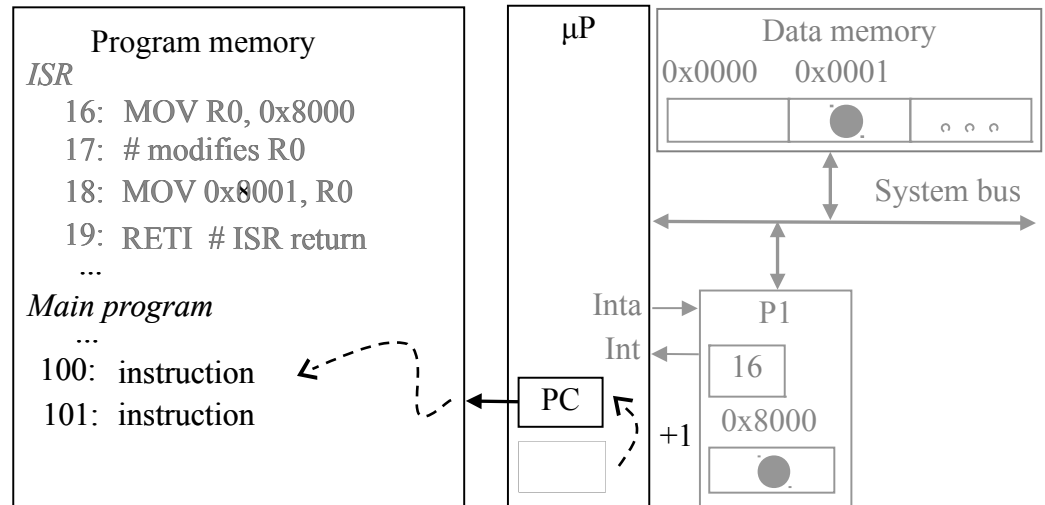
5(a): μ P jumps to the address on the bus (16). The ISR there reads data from 0x8000 and then writes it to 0x0001, which is in memory.

5(b): After being read, P1 de-asserts *Int*.



Peripheral to memory transfer *without* DMA, using vectored interrupt (cont')

6: The ISR returns, thus restoring PC to 100+1=101, where μ P resumes executing.



Additional interrupt issues

- Maskable vs. non-maskable interrupts
 - Maskable: programmer can set bit that causes processor to ignore interrupt
 - Important when in the middle of time-critical code
 - Non-maskable: a separate interrupt pin that can't be masked
 - Typically reserved for drastic situations, like power failure requiring immediate backup of data to non-volatile memory
- Jump to ISR
 - Some microprocessors treat jump same as call of any subroutine
 - Complete state saved (PC, registers) – may take hundreds of cycles
 - Others only save partial state, like PC only
 - Thus, ISR must not modify registers, or else must save them first
 - Assembly-language programmer must be aware of which registers stored

Interrupt-driven I/O, cont.

- Advantage over programmed I/O
 - Instead of waiting the operation to be finished, the CPU can do some useful work
- Still a problem
 - The CPU has to take care of each data unit (word), to move it to/from memory, and to issue an I/O command.
 - If large amounts of data have to be moved, this technique is still not efficient, because
 - the CPU has to take care for each data unit separately;
 - handling the interrupt also takes some time.

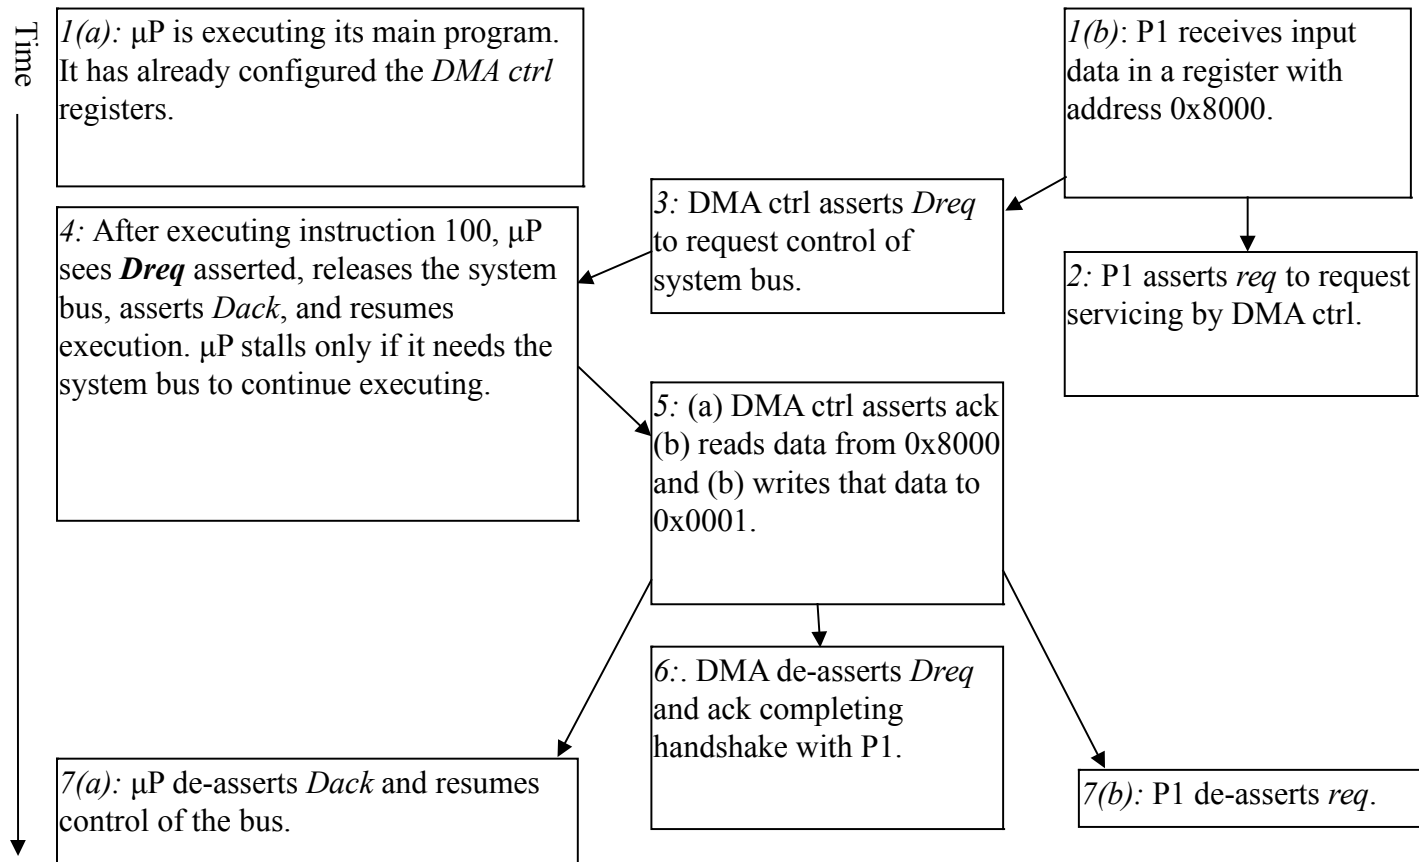
I/O Processing

1. Programmed I/O
2. Interrupt-driven I/O
3. **Direct memory access**

Direct memory access

- Buffering
 - Temporarily storing data in memory before processing
 - Data accumulated in peripherals commonly buffered
- Microprocessor could handle this with ISR
 - Storing and restoring microprocessor state inefficient
 - Regular program must wait
- DMA controller more efficient
 - Separate single-purpose processor
 - Microprocessor relinquishes control of system bus to DMA controller
 - Microprocessor can meanwhile execute its regular program
 - No inefficient storing and restoring state due to ISR call
 - Regular program need not wait unless it requires the system bus
 - Harvard architecture – processor can fetch and execute instructions as long as they don't access data memory – if they do, processor stalls

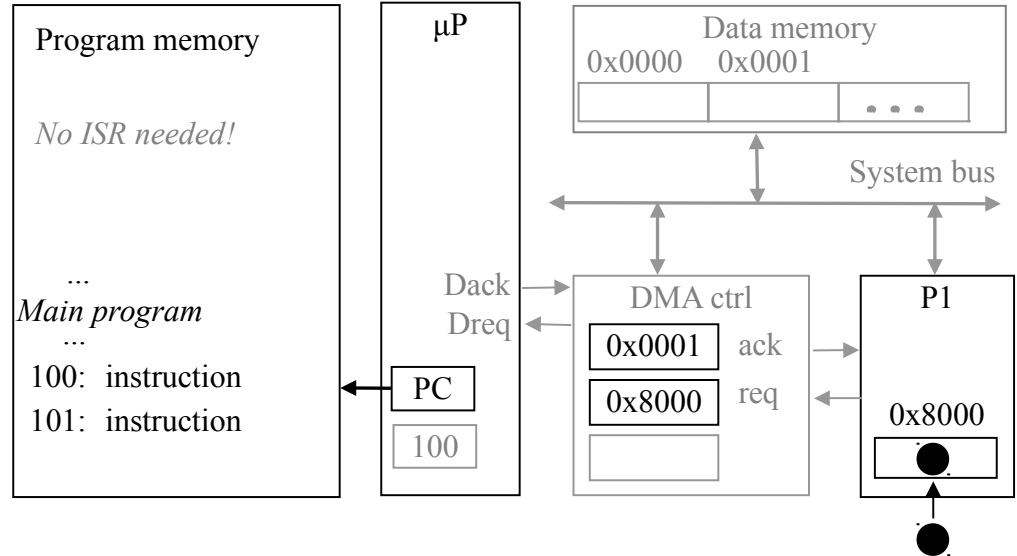
Peripheral to memory transfer with DMA



Peripheral to memory transfer with DMA (cont')

1(a): μ P is executing its main program. It has already configured the DMA ctrl registers

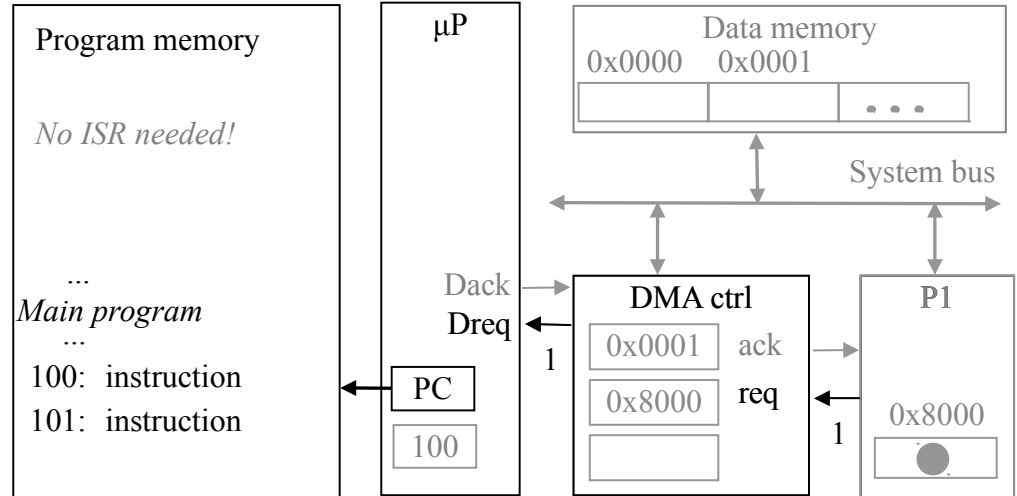
1(b): P1 receives input data in a register with address 0x8000.



Peripheral to memory transfer with DMA (cont')

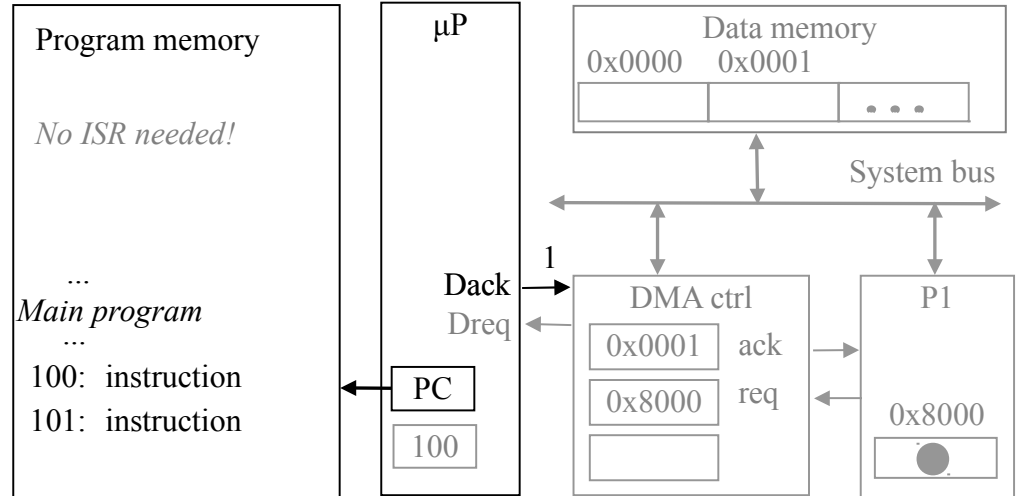
2: P1 asserts *req* to request servicing by DMA ctrl.

3: DMA ctrl asserts *Dreq* to request control of system bus



Peripheral to memory transfer with DMA (cont')

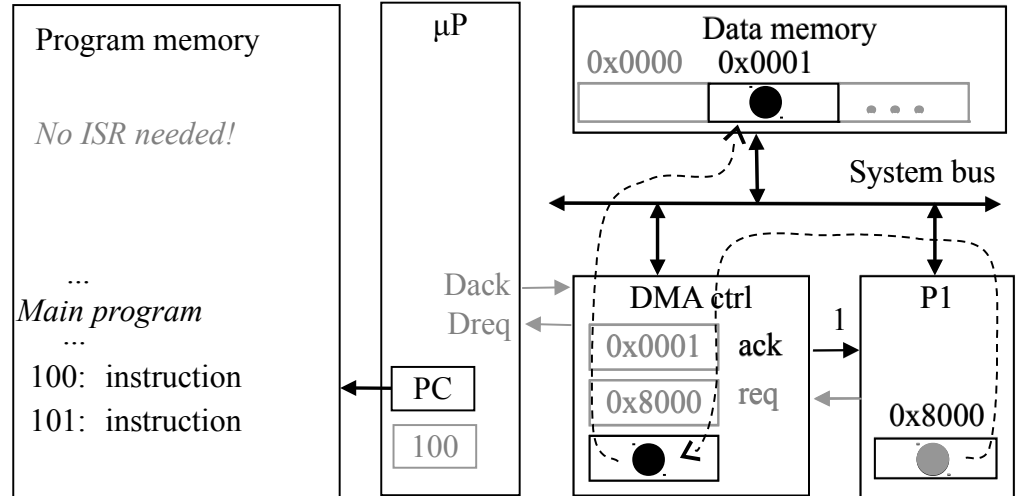
4: After executing instruction 100, μP sees *Dreq* asserted, releases the system bus, asserts *Dack*, and resumes execution, μP stalls only if it needs the system bus to continue executing.



Peripheral to memory transfer with DMA (cont')

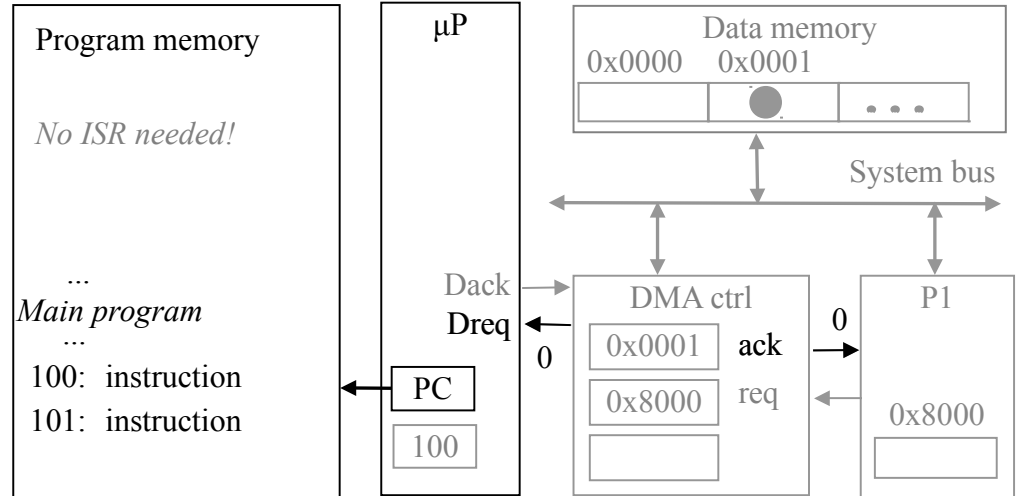
5: DMA ctrl (a) asserts ack, (b) reads data from 0x8000, and (c) writes that data to 0x0001.

(Meanwhile, processor still executing if not stalled!)



Peripheral to memory transfer with DMA (cont')

6: DMA de-asserts *Dreq* and *ack* completing the handshake with P1.



Analog conversion, timers and pulse-width modulation

Analog-to-digital converters


- Embedded systems interact with the environment
 - Embedded system understands **digital signals**
 - Discretely valued signals, such as integers, encoded in binary
 - Environment produces **analog signals**
 - Continuous valued signal, such as temperature or speed, with infinite possible values
- Analog to digital converter (ADC, A/D, A2D)
 - Converts an analog signal to a digital signal
- Digital to analog converter (DAC, D/A, D2A)
 - Converts a digital signal to analog signal

Analog-to-digital converters


$V_{\max} = 7.5V$

7.5V	1111
7.0V	1110
6.5V	1101
6.0V	1100
5.5V	1011
5.0V	1010
4.5V	1001
4.0V	1000
3.5V	0111
3.0V	0110
2.5V	0101
2.0V	0100
1.5V	0011
1.0V	0010
0.5V	0001
0V	0000

proportionality




analog to digital



digital to analog


Timers, counters, watchdog timers

- **Timer: measures time intervals**
 - To generate timed output events
 - e.g., hold traffic light green for 10 s
 - To measure input events
 - e.g., measure a car's speed
- **Based on counting clock pulses**
 - Range
 - Maximum measurable time interval
 - Resolution
 - Minimum measurable time interval
 - Example:
 - E.g., let Clk period be 10 ns
 - And we count 20,000 Clk pulses
 - Then 200 microseconds have passed
 - 16-bit counter would count up to $65,535 \times 10 \text{ ns} = 655.35 \text{ microsec.}$, resolution = 10 ns
 - Top: indicates top count reached, wrap-around




Counters

- Counter: like a timer, but counts pulses on a general input signal rather than clock
 - e.g., count cars passing over a sensor
 - Can often configure device as either a timer or counter



Example: Reaction Timer



- Measure time between turning light on and user pushing button

```
/* main.c */

#define MS_INIT 63535
void main(void){
    int count_milliseconds = 0;

    configure timer mode
    set Cnt to MS_INIT

    wait a random amount of time
    turn on indicator light
    start timer

    while (user has not pushed reaction button){
        if(Top) {
            stop timer
            set Cnt to MS_INIT
            start timer
            reset Top
            count_milliseconds++;
        }
    }
    turn light off
    printf("time: %i ms", count_milliseconds);
}
```

Watchdog timer

- Special kind of timer
 - Must reset timer every X time unit, otherwise the watchdog timer generates a signal
 - The watchdog timer “times-out” if we fail to reset it

Watchdog timer

- Common use: detect failure, self-reset
- Another use: timeouts
- Example: ATM machine
 - Time-out after two minutes if no key pressed

Figure 1: A typical watchdog setup



Watchdog timer : Code

```
uint16 volatile * pWatchdog = (uint16 volatile *) 0xFF0000;  
main(void) {  
    hwinit();  
    for (;;) {  
        *pWatchdog = 10000;  
        read_sensors();  
        control_motor();  
        display_status();  
    }  
}
```

Watchdog timer : Software anomalies

- Logical fallacy, e.g., infinite loop
 - What if this happens with `read_sensors()`?
- Unusually high number of interrupts arrives – leads to dangerous delay in feeding the motor new control instructions
- Deadlocks that occur in multi-tasking kernels
- Note: failed hardware can lead to constant resets. So it is a good idea to count resets & give up after fixed number of failures

Pulse width modulator (PWM)

- PWM : a powerful technique to control analog circuits with a microprocessor's digital output
- Essentially, it generates pulses with specific high/low times
- Duty cycle: % time high
 - Square wave: 50% duty cycle

Simple PWM circuit

- Lamp connected to 9 V battery
- Switch on for 50 ms and then switch off for 50ms
- The lamp behaves as if it was powered by 4.5V source
- Duty cycle 50%
- Switching off and on frequency should be much higher than response time of the system



Controlling a DC motor with a PWM


- DC motor
 - Speed of DC motor is proportional to the average voltage applied
 - We set the duty cycle of a PWM such that we obtain the desired average voltage, see examples

50% Duty cycle

- Example: create 2.5V signal given a digital source of 5V, what should be the duty cycle of the PWM?

50% Duty cycle

- Example: create 2.5V signal given a digital source of 5V, what should be the duty cycle of the PWM?




50% duty cycle – average pwm_o is 2.5V.

25% Duty cycle

- Example: create 1.25V signal given a digital source of 5V, what should be the duty cycle of the PWM?

25% Duty cycle

- Example: create 1.25V signal given a digital source of 5V, what should be the duty cycle of the PWM?



25% duty cycle – average pwm_o is 1.25V

Extra material

Microcontroller, cont.

- For embedded control applications
 - Reading sensors, setting actuators
 - Mostly dealing with events (bits): data is present, but not in huge amounts
 - e.g., automotive control systems, wearable health-care devices, pacemaker, disk drive, digital camera, washing machine, microwave oven
- Microcontroller features
 - On-chip peripherals
 - Timers, analog-digital converters, serial communication, etc.
 - Tightly integrated for programmer, typically part of register space
 - On-chip program and data memory
 - Direct programmer access to many of the chip's pins
 - Specialized instructions for bit-manipulation and other low-level operations

Direct Memory Access (DMA)

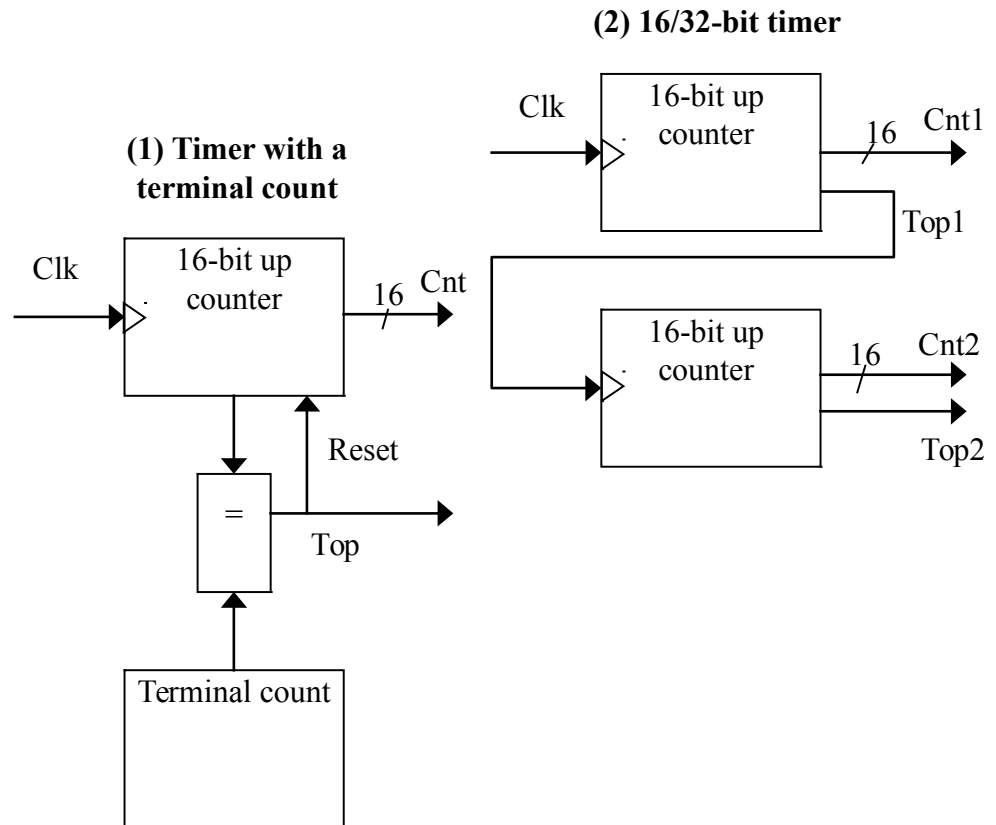
- An additional module on the system bus, the DMA module, takes care of the I/O transfer for the whole sequence of data.
 - The CPU issues a command to the DMA module and transfers to it all the needed information.
 - When starting an operation, the CPU informs the DMA module about:
 - what operation (read or write);
 - the address of the I/O device involved;
 - the starting location in memory where information has to be stored to or read from;
 - the number of words to be transferred.
 - The DMA module performs all the operations (acting similarly to the CPU with programmed I/O): it transfers all the data between I/O module and memory without going through the CPU.
 - When the DMA module has finished, it issues an interrupt to the CPU.

Other timer structures

1. Interval timer


- Indicates when desired time interval has passed
- We set terminal count to desired interval
 - $\text{Number of clock cycles} = \text{Desired time interval} / \text{Clock period}$

2. Cascaded counters



75% Duty cycle

- Example: create 3.75V signal given a digital source of 5V, what should be the duty cycle of the PWM?



75% duty cycle – average pwm_o is 3.75V.