
Concurrency

Lecture outline

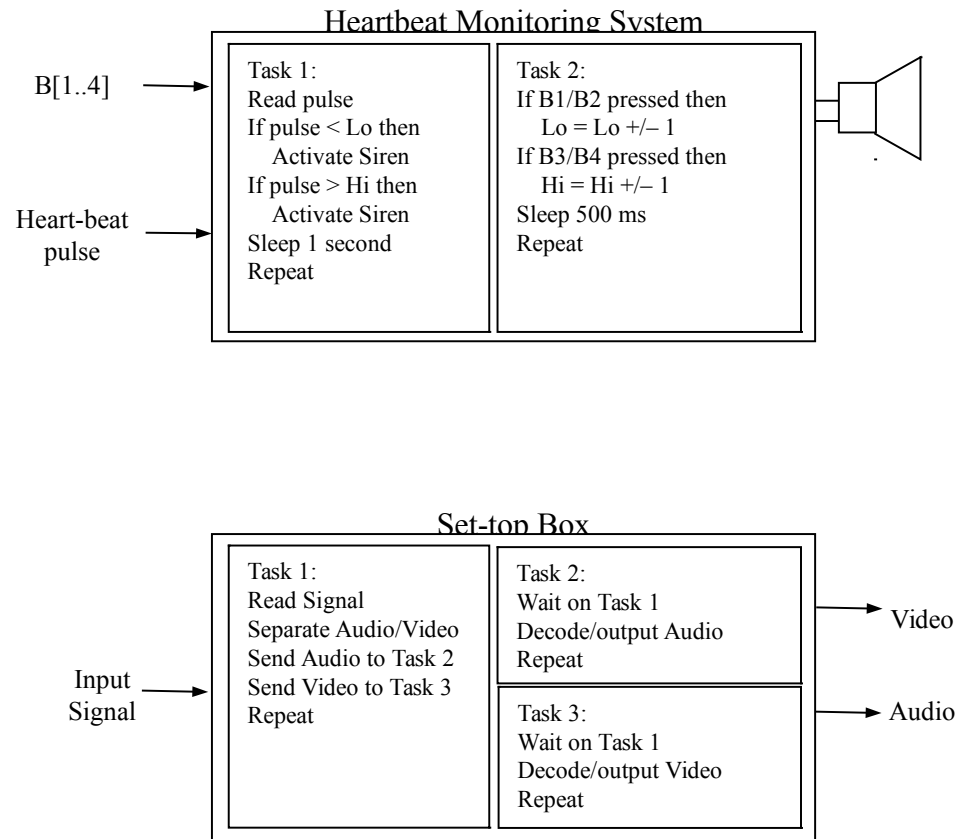
- Why concurrency?
- Foreground/background vs. multi-tasking systems
- Concurrent processes
 - Communication: message passing vs. shared memory
- Scheduling
- Bus scheduling

Lecture outline


- **Why concurrency?**
- Foreground/background vs. multi-tasking systems
- Concurrent processes
 - Communication: message passing vs. shared memory
- Scheduling
- Bus scheduling

Why concurrency?


- Consider two examples having separate tasks running independently but sharing data
- Difficult to write system using sequential program model
- Concurrent process model easier
 - Separate sequential programs (processes) for each task
 - Programs communicate with each other



Concurrent Programs



A **sequential** program has a single thread of control.



A **concurrent** program has multiple threads of control allowing it perform multiple computations “in parallel” and to control multiple external activities which occur at the same time.

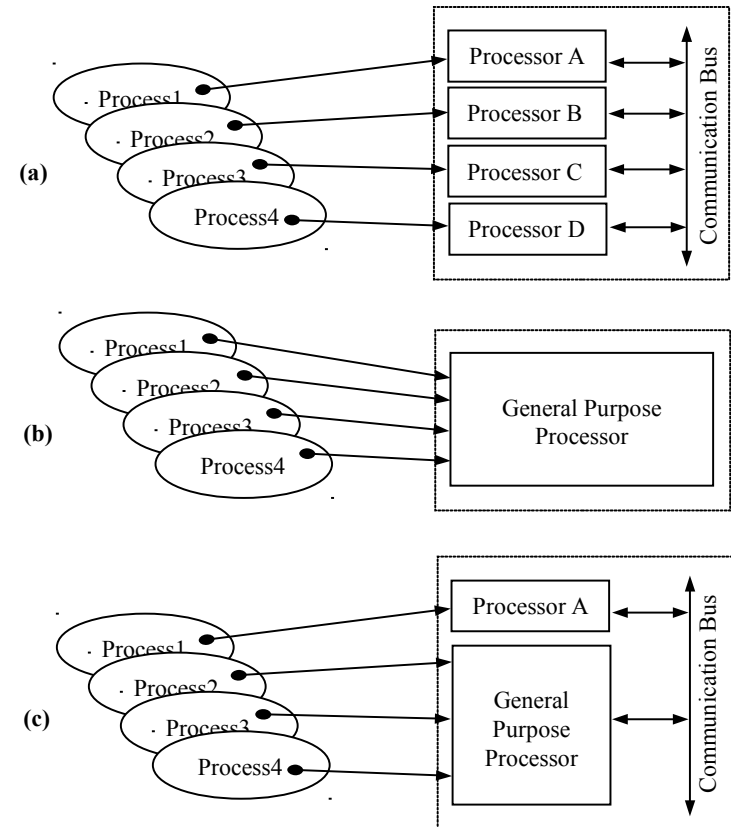
[Magee and Kramer 2006]

Concurrency manifestations

- Multiple applications
 - Multiprogramming
- Structured application
 - Application can be a set of concurrent processes
- Operating-system structure
 - Operating system is a set of processes or threads

Concurrent process model: implementation

- Can use single and/or general-purpose processors
- (a) Multiple processors, each executing one process
 - True multitasking (parallel processing)
 - General-purpose processors
 - Use programming language like C and compile to instructions of processor
 - Expensive and in most cases not necessary
 - Custom single-purpose processors
 - More common
- (b) One general-purpose processor running all processes
 - Most processes don't use 100% of processor time
 - Can share processor time and still achieve necessary execution rates
- (c) Combination of (a) and (b)
 - Multiple processes run on one general-purpose processor while one or more processes run on own single-purpose processor



Challenges with concurrency

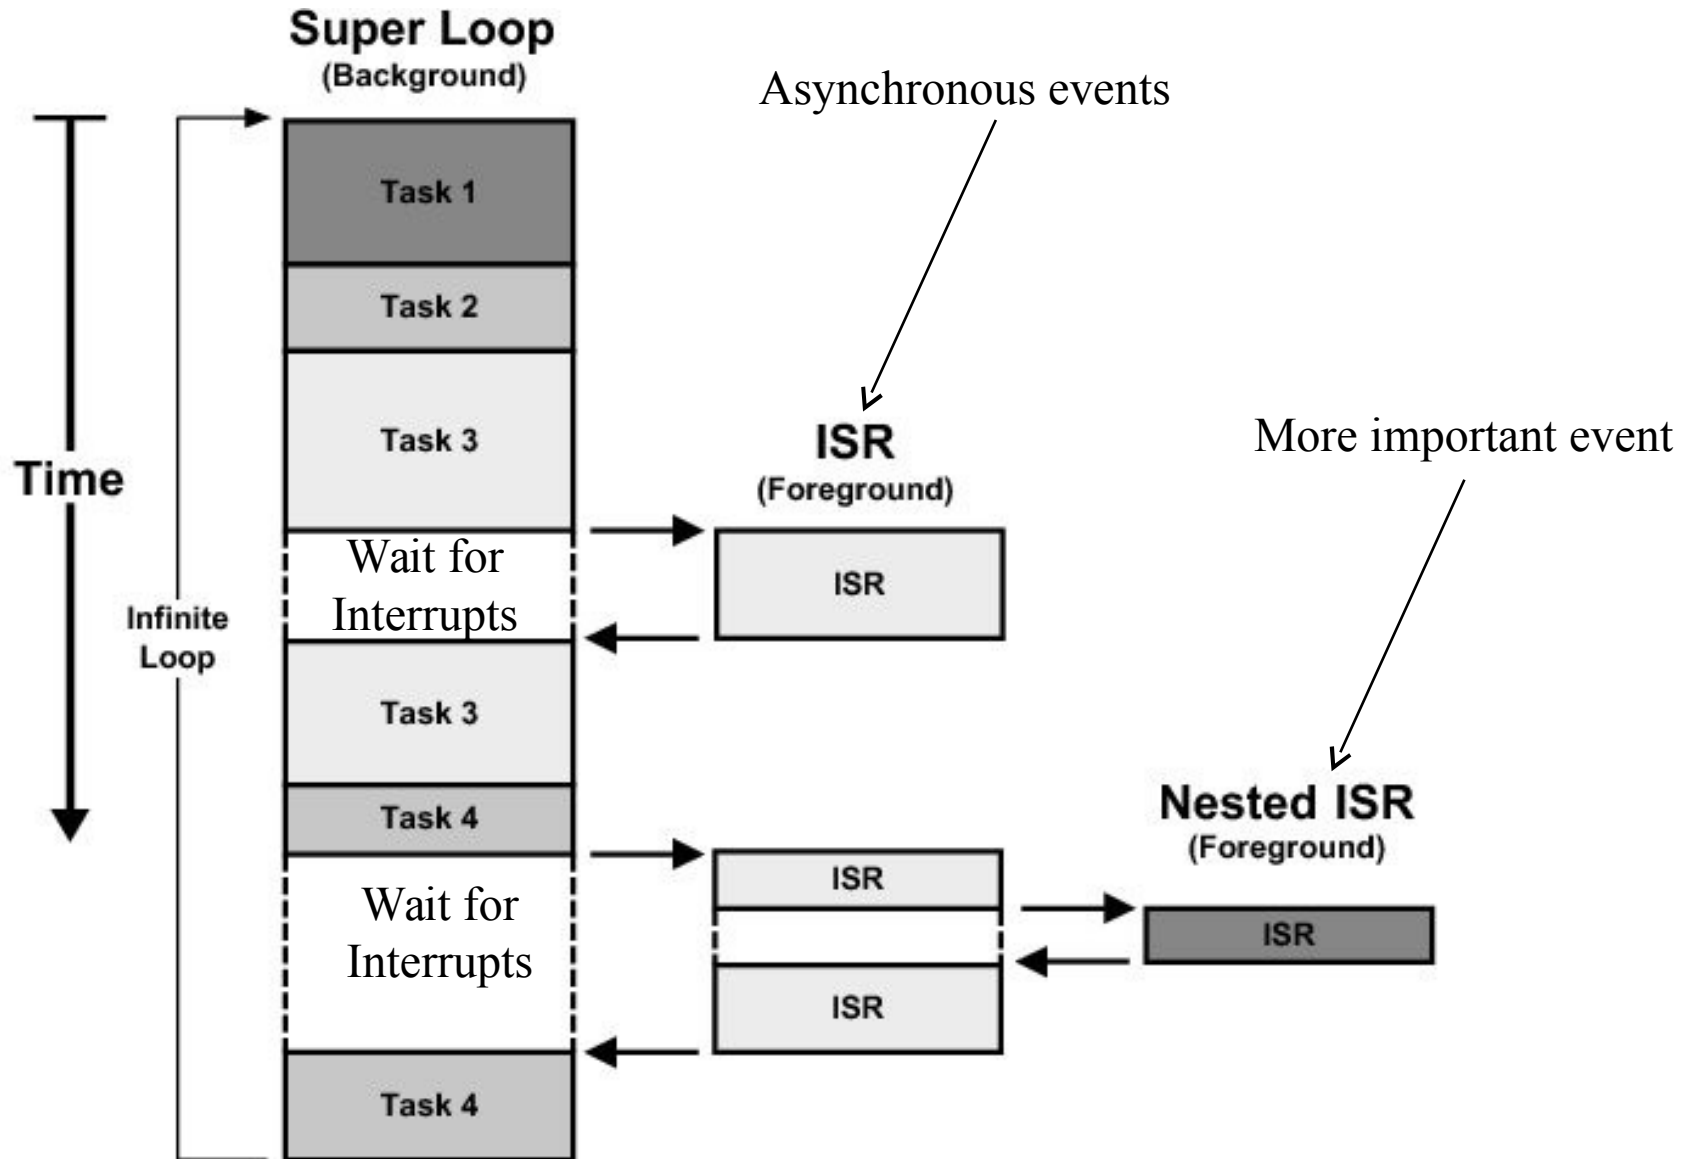
- Sharing global resources
- Management of allocation of resources
- Programming errors difficult to locate

Lecture outline

- Why concurrency?
- **Foreground/background vs. multi-tasking systems**
- Concurrent processes
 - Communication: message passing vs. shared memory
- Scheduling
- Bus scheduling

Foreground/Background systems

- Efficient for small systems of low complexity
- Infinite loop that call modules or tasks to perform the desired operations (also called task level or superloop)
- Interrupt Service Routines (ISRs) handle asynchronous events (foreground also called ISR level)
 - Timer interrupts
 - I/O interrupts



Foreground/Background systems

- Critical tasks are handled by ISRs to ensure they perform in timely fashion
- Information for a background module that makes an ISR available is not processed until the background routine gets its turn to execute. This is called task-level response.
- The worst-case task-level response is depends on how long the background loop takes to execute and since the execution time is not constant, this is difficult to predict.
- High volume and low-cost microcontroller-based applications (e.g., microwaves, telephones,...) are designed are foreground/background systems

Foreground/Background

```
/* Background */
void main (void)
{
    Initialization;
    FOREVER {
        Read analog inputs;
        Read discrete inputs;
        Perform monitoring functions;
        Perform control functions;
        Update analog outputs;
        Update discrete outputs;
        Scan keyboard;
        Handle user interface;
        Update display;
        Handle communication requests;
        Other...
    }
}

/* Foreground */
ISR (void)
{
    Handle asynchronous event;
}
```

Foreground/Background: Advantages

- Used in low cost embedded applications
- Memory requirements only depends on your application
- Single stack area for:
 - Function nesting
 - Local variables
 - ISR nesting
- Minimal interrupt latency for bare minimum embedded systems

Foreground/Background: Disadvantages

- Background response time is the background execution time
 - Non-deterministic, affected by if, for, while ...
 - May not be responsive enough
 - Changes as you change your code

Foreground/Background: Disadvantages

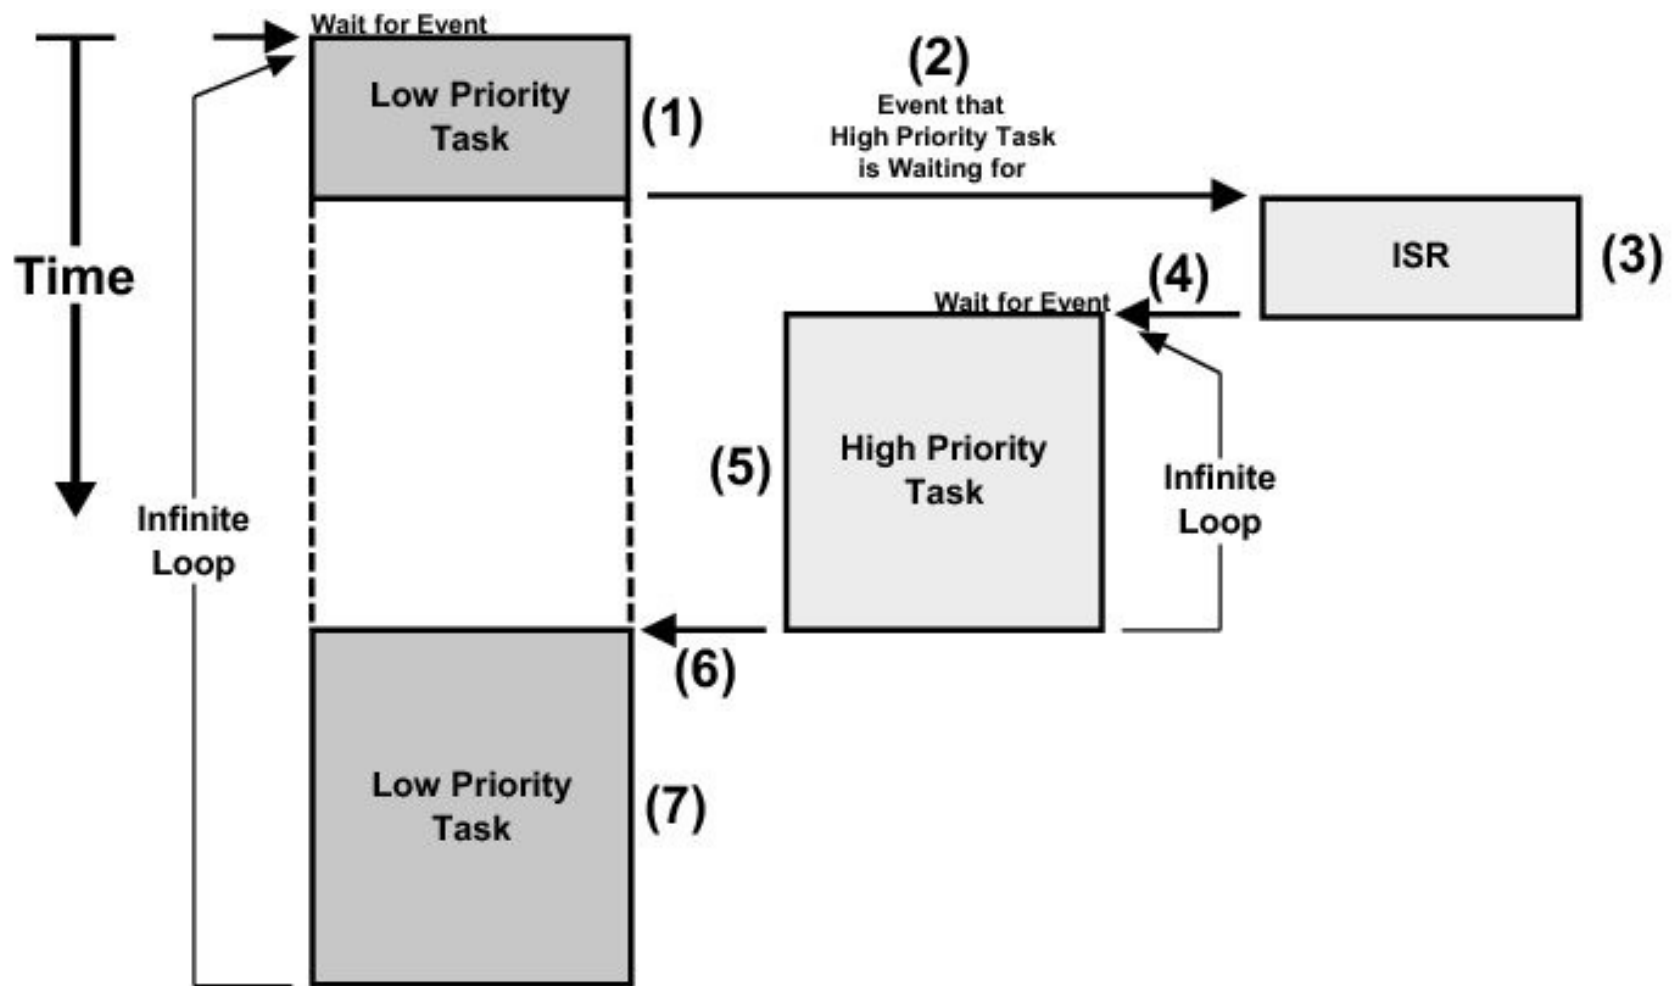
- All 'tasks' have the same priority!
 - Code executes in sequence
 - If an important event occurs it's handled at the same priority as everything else!
 - You may need to execute the same code often to avoid missing an event.

Foreground/Background: Disadvantages

- Code is harder to maintain and can become messy
 - Imagine the C program as the number of tasks increase!

Migrate to multi-tasking systems

- Each operation in the superloop/background is broken apart into a task, that by itself runs in infinite loop



Multi-tasking system

- Each task is a simple program that thinks it has the entire CPU to itself, and typically executed in an infinite loop.
- In the CPU only one task runs at any given time. This is management --- scheduling and switching the CPU between several tasks --- is performed by the kernel of the real-time system

Lecture outline

- Why concurrency?
- Foreground/background vs. multi-tasking systems
- **Concurrent processes**
 - Communication: message passing vs. shared memory
- Scheduling
- Bus scheduling

Concurrent processes

- Only one process runs at a time while others are *suspended*.
- Processor switches from one process to another so quickly that it appears all processes are running simultaneously. Processes run *concurrently*.
- Programmer assigns *priority* to each process and the *scheduler* uses this to determine which process to run next.

Process


- A sequential program, typically an infinite loop
 - Executes concurrently with other processes
- Basic operations on processes
 - Create and terminate
 - Create is like a procedure call but caller doesn't wait
 - Created process can itself create new processes
 - Terminate kills a process, destroying all data
 - Suspend and resume
 - Suspend puts a process on hold, saving state for later execution
 - Resume starts the process again where it left off
 - Join
 - A process suspends until a particular child process finishes execution

Processes vs. threads

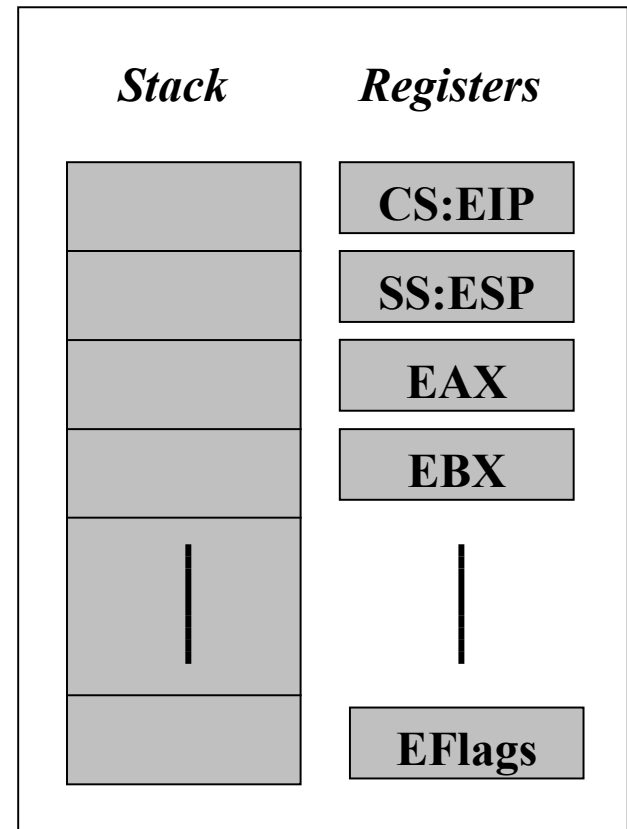
- Different meanings in operating system terminology
- Regular processes
 - Heavyweight process
 - Own virtual address space (stack, data, code)
 - System resources (e.g., open files)
- Threads
 - Lightweight process
 - Subprocess within process
 - Only program counter, stack, and registers
 - Shares address space, system resources with other threads
 - Allows quicker communication between threads
 - Small compared to heavyweight processes
 - Can be created quickly
 - Low cost switching between threads

Each process maintains its own stack and register contents

Context of Process 1



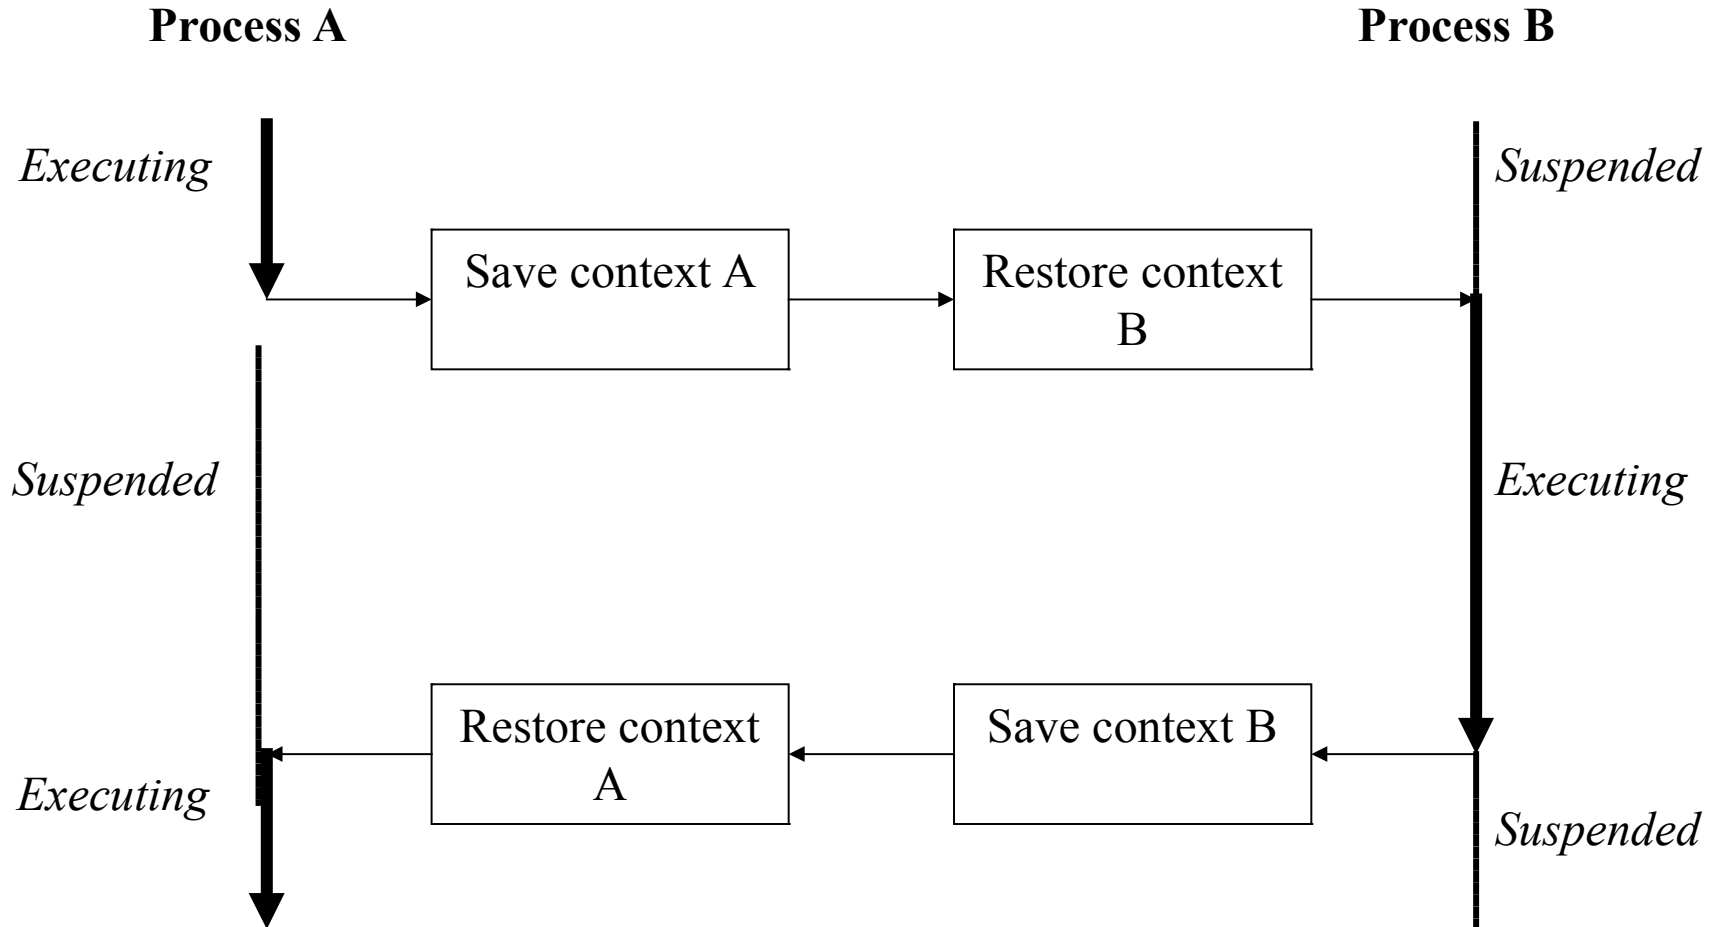
Context of Process N



Context switching

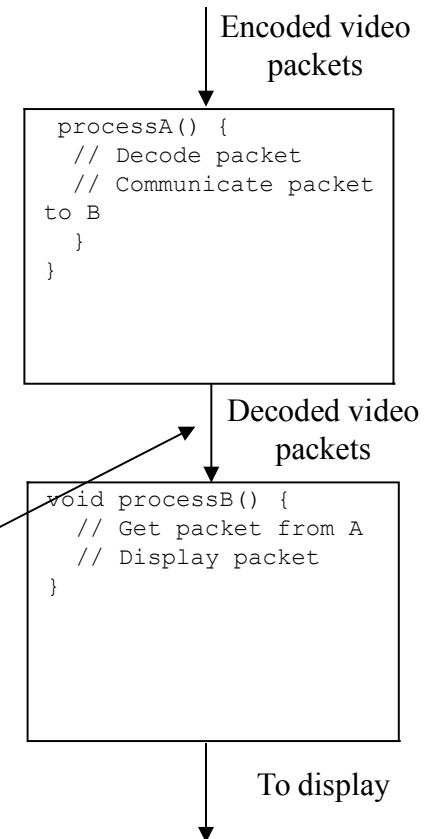
- Each process has its own stack and a special region of memory referred to as its *context*.
- A *context switch* from process "A" to process "B" first saves all CPU registers in context A, and then reloads all CPU registers from context B.
- Since CPU registers includes SS:ESP and CS:EIP, reloading context B reactivates process B's stack and returns to where it left off when it was last suspended.

Context switching, cont.



Communication among processes

- Processes need to communicate data and signals to solve their computation problem
 - Processes that don't communicate are just independent programs solving separate problems
- Basic example: producer/consumer
 - Process A produces data items, Process B consumes them
 - E.g., A decodes video packets, B display decoded packets on a screen
- How do we achieve this communication?
 - Two basic methods
 - Shared memory
 - Message passing



Shared Memory

- Processes read and write shared variables
 - No time overhead, easy to implement
 - But, hard to use – mistakes are common
- Example: Producer/consumer with a mistake
 - Share *buffer[N]*, *count*
 - *count* = # of valid data items in *buffer*
 - *processA* produces data items and stores in *buffer*
 - If *buffer* is full, must wait
 - *processB* consumes data items from *buffer*
 - If *buffer* is empty, must wait
 - Error when both processes try to update *count* concurrently (lines 10 and 19) and the following execution sequence occurs. Say “*count*” is 3.
 - A loads *count* (*count* = 3) from memory into register R1 (R1 = 3)
 - A increments R1 (R1 = 4)
 - B loads *count* (*count* = 3) from memory into register R2 (R2 = 3)
 - B decrements R2 (R2 = 2)
 - A stores R1 back to *count* in memory (*count* = 4)
 - B stores R2 back to *count* in memory (*count* = 2)
 - *count* now has incorrect value of 2

```
01: data_type buffer[N];
02: int count = 0;
03: void processA() {
04:     int i;
05:     while( 1 ) {
06:         produce(&data);
07:         while( count == N ); /*loop*/
08:         buffer[i] = data;
09:         i = (i + 1) % N;
10:         count = count + 1;
11:     }
12: }
13: void processB() {
14:     int i;
15:     while( 1 ) {
16:         while( count == 0 ); /*loop*/
17:         data = buffer[i];
18:         i = (i + 1) % N;
19:         count = count - 1;
20:         consume(&data);
21:     }
22: }
23: void main() {
24:     create_process(processA);
25:     create_process(processB);
26: }
```

Message Passing

- Message passing
 - Data explicitly sent from one process to another
 - Sending process performs special operation, *send*
 - Receiving process must perform special operation, *receive*, to receive the data
 - Both operations must explicitly specify which process it is sending to or receiving from
 - Safer model, but less flexible

```
void processA() {  
    while( 1 ) {  
        produce(&data)  
        send(B, &data);  
        /* region 1 */  
        receive(B, &data);  
        consume(&data);  
    }  
}
```

```
void processB() {  
    while( 1 ) {  
        receive(A, &data);  
        transform(&data)  
        send(A, &data);  
        /* region 2 */  
    }  
}
```

Back to Shared Memory: Mutual Exclusion

- Certain sections of code should not be performed concurrently
 - Critical section
 - section of code where simultaneous updates, by multiple processes to a shared memory location, can occur
- When a process enters the critical section, all other processes must be locked out until it leaves the critical section
 - Mutex
 - A shared object used for locking and unlocking segment of shared data
 - Disallows read/write access to memory it guards
 - Multiple processes can perform lock operation simultaneously, but only one process will acquire lock
 - All other processes trying to obtain lock will be put in blocked state until unlock operation performed by acquiring process when it exits critical section
 - These processes will then be placed in runnable state and will compete for lock again

Correct Shared Memory Solution to the Consumer-Producer Problem

- The primitive *mutex* is used to ensure critical sections are executed in mutual exclusion of each other
- Following the same execution sequence as before:
 - A/B execute *lock* operation on *count_mutex*
 - Either A or B will acquire *lock*
 - Say B acquires it
 - A will be put in blocked state
 - B loads *count* (*count* = 3) from memory into register R2 (R2 = 3)
 - B decrements R2 (R2 = 2)
 - B stores R2 back to *count* in memory (*count* = 2)
 - B executes *unlock* operation
 - A is placed in runnable state again
 - A loads *count* (*count* = 2) from memory into register R1 (R1 = 2)
 - A increments R1 (R1 = 3)
 - A stores R1 back to *count* in memory (*count* = 3)
- *Count* now has correct value of 3

```
01: data_type buffer[N];
02: int count = 0;
03: mutex count_mutex;
04: void processA() {
05:     int i;
06:     while( 1 ) {
07:         produce(&data);
08:         while( count == N );/*loop*/
09:         buffer[i] = data;
10:         i = (i + 1) % N;
11:         count_mutex.lock();
12:         count = count + 1;
13:         count_mutex.unlock();
14:     }
15: }
16: void processB() {
17:     int i;
18:     while( 1 ) {
19:         while( count == 0 );/*loop*/
20:         data = buffer[i];
21:         i = (i + 1) % N;
22:         count_mutex.lock();
23:         count = count - 1;
24:         count_mutex.unlock();
25:         consume(&data);
26:     }
27: }
28: void main() {
29:     create_process(processA);
30:     create_process(processB);
31: }
```


A common problem in concurrent programming: deadlock

- Deadlock: A condition where 2 or more processes are blocked waiting for the other to unlock critical sections of code
 - Both processes are then in blocked state
 - Cannot execute unlock operation so will wait forever
- Example code has 2 different critical sections of code that can be accessed simultaneously
 - 2 locks needed (mutex1, mutex2)
 - Following execution sequence produces deadlock
 - A executes lock operation on *mutex1* (and acquires it)
 - B executes lock operation on *mutex2* (and acquires it)
 - A/B both execute in critical sections 1 and 2, respectively
 - A executes lock operation on *mutex2*
 - A blocked until B unlocks *mutex2*
 - B executes lock operation on *mutex1*
 - B blocked until A unlocks *mutex1*
 - DEADLOCK!
- One deadlock elimination protocol requires locking of numbered mutexes in increasing order and two-phase locking (2PL)
 - Acquire locks in 1st phase only, release locks in 2nd phase

```
01: mutex mutex1, mutex2;
02: void processA() {
03:     while( 1 ) {
04:         ...
05:         mutex1.lock();
06:         /* critical section 1 */
07:         mutex2.lock();
08:         /* critical section 2 */
09:         mutex2.unlock();
10:         /* critical section 1 */
11:         mutex1.unlock();
12:     }
13: }
14: void processB() {
15:     while( 1 ) {
16:         ...
17:         mutex2.lock();
18:         /* critical section 2 */
19:         mutex1.lock();
20:         /* critical section 1 */
21:         mutex1.unlock();
22:         /* critical section 2 */
23:         mutex2.unlock();
24:     }
25: }
```

Summary:

multiple processes sharing single processor

- Manually rewrite processes as a single sequential program
 - Ok for simple examples, but extremely difficult for complex examples
 - Automated techniques have evolved but not common
- Can use multitasking operating system
 - Much more common
 - Operating system schedules processes, allocates storage, and interfaces to peripherals, etc.
 - Real-time operating system (RTOS) can guarantee execution rate constraints are met
 - Describe concurrent processes with languages having built-in processes (Java, Ada, etc.) or a sequential programming language with library support for concurrent processes (C, C++, etc. using POSIX threads for example)
- Can convert processes to sequential program with process scheduling right in code
 - Less overhead (no operating system)
 - More complex/harder to maintain

Lecture outline

- Why concurrency?
- Foreground/background vs. multi-tasking systems
- Concurrent processes
 - Communication: message passing vs. shared memory
- **Scheduling**
- Bus scheduling

Implementation: process scheduling

- Must meet timing requirements when multiple concurrent processes implemented on single general-purpose processor
- Scheduler
 - Special process that decides when and for how long each process is executed
 - Implemented as preemptive or nonpreemptive scheduler

Preemptive vs non-preemptive

- Preemptive
 - Determines how long a process executes before preempting to allow another process to execute
 - Time quantum: predetermined amount of execution time preemptive scheduler allows each process (may be 10 to 100s of milliseconds long)
 - Determines which process will be next to run
- Nonpreemptive
 - Only determines which process is next after current process finishes execution

Static vs dynamic scheduling

- Static (off-line)
 - complete a priori knowledge of the task set and its constraints is available
 - hard/safety-critical system
- Dynamic (on-line)
 - partial taskset knowledge, runtime predictions
 - firm/soft/best-effort systems, hybrid systems

Scheduling Approaches

- Cyclic executives
- Fixed priority scheduling
 - RM - Rate Monotonic
 - DM - Deadline Monotonic Scheduling
- Dynamic priority scheduling
 - EDF - Earliest Deadline First
 - LSF - Least Slack First

Cyclic Executive

Process	Period	Comp. Time
A	25	10
B	25	8
C	50	5
D	50	4
E	100	2

```


loop
  Wait_For_Interrupt;
  Procedure_For_A;
  Procedure_For_B;
  Procedure_For_C;

  Wait_For_Interrupt;
  Procedure_For_A;
  Procedure_For_B;
  Procedure_For_D;
  Procedure_For_E;

  Wait_For_Interrupt;
  Procedure_For_A;
  Procedure_For_B;
  Procedure_For_C;

  Wait_For_Interrupt;
  Procedure_For_A;
  Procedure_For_B;
  Procedure_For_D;

end loop;
  
```

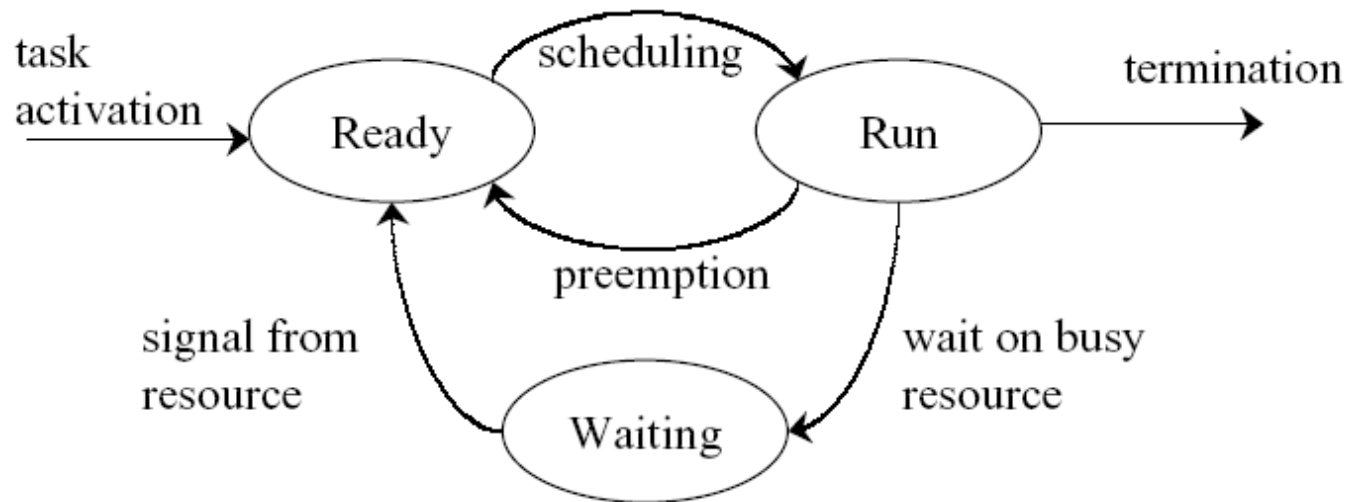


- What happens if the periods are not harmonic?

Priority-based scheduling

- Every task has an associated priority
- Run task with the highest priority
 - At every scheduling decision moment
- Examples
 - Rate Monotonic (RM)
 - Static priority assignment
 - Earliest Deadline First (EDF)
 - Dynamic priority assignment
 - And many others ...

States of a process



Schedulability Test

- Test to determine whether a feasible schedule exists
- Sufficient
 - + if test is passed, then tasks are definitely schedulable
 - if test is not passed, we don't know
- Necessary
 - + if test is passed, we don't know
 - if test is not passed, tasks are definitely not schedulable
- Exact
 - sufficient & necessary at the same time

Rate Monotonic

- Each process is assigned a (unique) priority based on its period; the shorter the period, the higher the priority
- Assumes the “Simple task model”


- Fixed priority scheduling

Process	Period	Priority
A	25	5
B	60	3
C	42	4
D	105	1
E	75	2

- Preemptive
 - Unless stated otherwise


Example 1

- Assume we have the following task set
 - OBS: not scheduled yet ...



Example 1 (cont'd)

- Scheduled with RM



Schedulability test for RM

- Sufficient, but not necessary:

$$\sum_{i=1}^N \left(\frac{C_i}{T_i} \right) \leq N (2^{1/N} - 1)$$

N	Utilization Bound
1	100.0%
2	82.8%
3	78.0%
4	75.7%
5	74.3%
10	71.8%

- Necessary, but not sufficient:

In the limit: 69.3%

$$\sum_{i=1}^N \left(\frac{C_i}{T_i} \right) \leq 1$$

Example 2


- Taskset
 - Period (T_i)
 - WCET (C_i)
- | | P1 | P2 | P3 |
|--|----|----|----|
| | 20 | 50 | 30 |
| | 7 | 10 | 5 |
- Is this schedulable?

Example 3

- Taskset

	Period	Comp. Time	Priority	Utilization
Task_1	80	40	1	0.50
Task_2	40	10	2	0.25
Task_3	20	5	3	0.25

- Gantt chart:



Exact schedulability test

- The schedulability of a given taskset for RM can be decided by:
 - Drawing a schedule
 - Doing a response time analysis
- Complexity: Pseudo-polynomial time

Optimality of scheduling algorithms

- “A scheduler is optimal if it always finds a schedule when a schedulability test indicates there is one.”
 - Burns, 1991
- “An optimal scheduling algorithm is one that may fail to meet a deadline if no other scheduling algorithm can meet it.”
 - Stankovic et al., 1995
- “An optimal scheduling algorithm is guaranteed to always find a feasible schedule, given that a feasible schedule does exist.”
 - Hansson, 1998

Optimality of RM

- Rate Monotonic is optimal among fixed priority schedulers
 - If we assume the “Simple Process Model” for the tasks

What to do if not schedulable

- Change the task set utilisation
 - by reducing C_i
 - code optimisation
 - faster processor
- Increase T_i for some process
 - If your program and environment allows it

RM characteristics

- Easy to implement.
- Drawback:
 - May not give a feasible schedule even if processor is idle at some points.

Earliest Deadline First (EDF)

- Always runs the process that is closest to its deadline.
- Dynamic priority scheduling
 - Evaluated at run-time
 - What are the events that should trigger a priority reevaluation?
- Assumes the “Simple task model”
 - Actually more relaxed: $D_i < T_i$
- Preemptive
 - Unless stated otherwise

Schedulability test for EDF


- Utilisation test
 - Necessary and sufficient (exact)

$$\sum_i^N \left(\frac{C_i}{T_i} \right) \leq 1$$

Optimality of EDF

- EDF is optimal among dynamic priority schedulers
 - If we assume the “Simple Process Model” for the tasks
 - Or a more relaxed one where $D_i < T_i$

Domino Effect

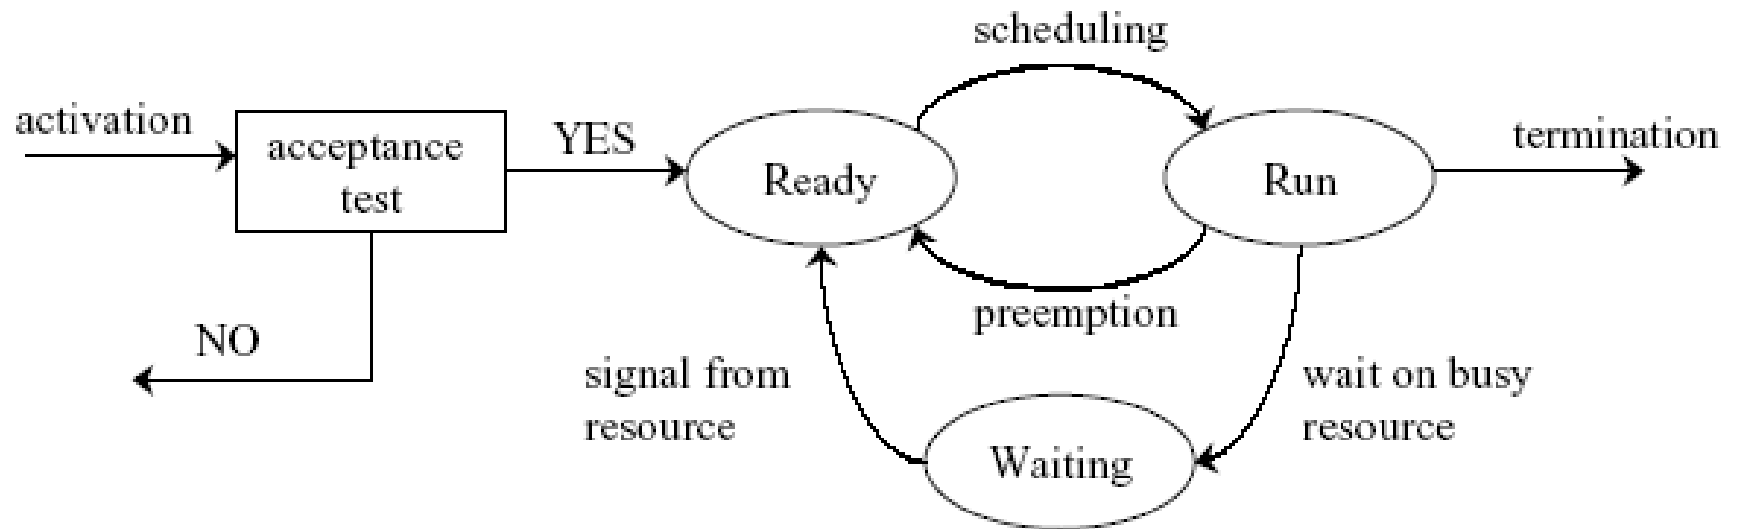


Domino effect!!!

EDF vs. RM

- EDF can handle tasksets with higher processor utilisation.
- EDF has simpler exact analysis
- RMS can be implemented to run faster at run-time
 - Depends on the OS
 - But they usually like fixed priorities more

Dynamic Scheduling



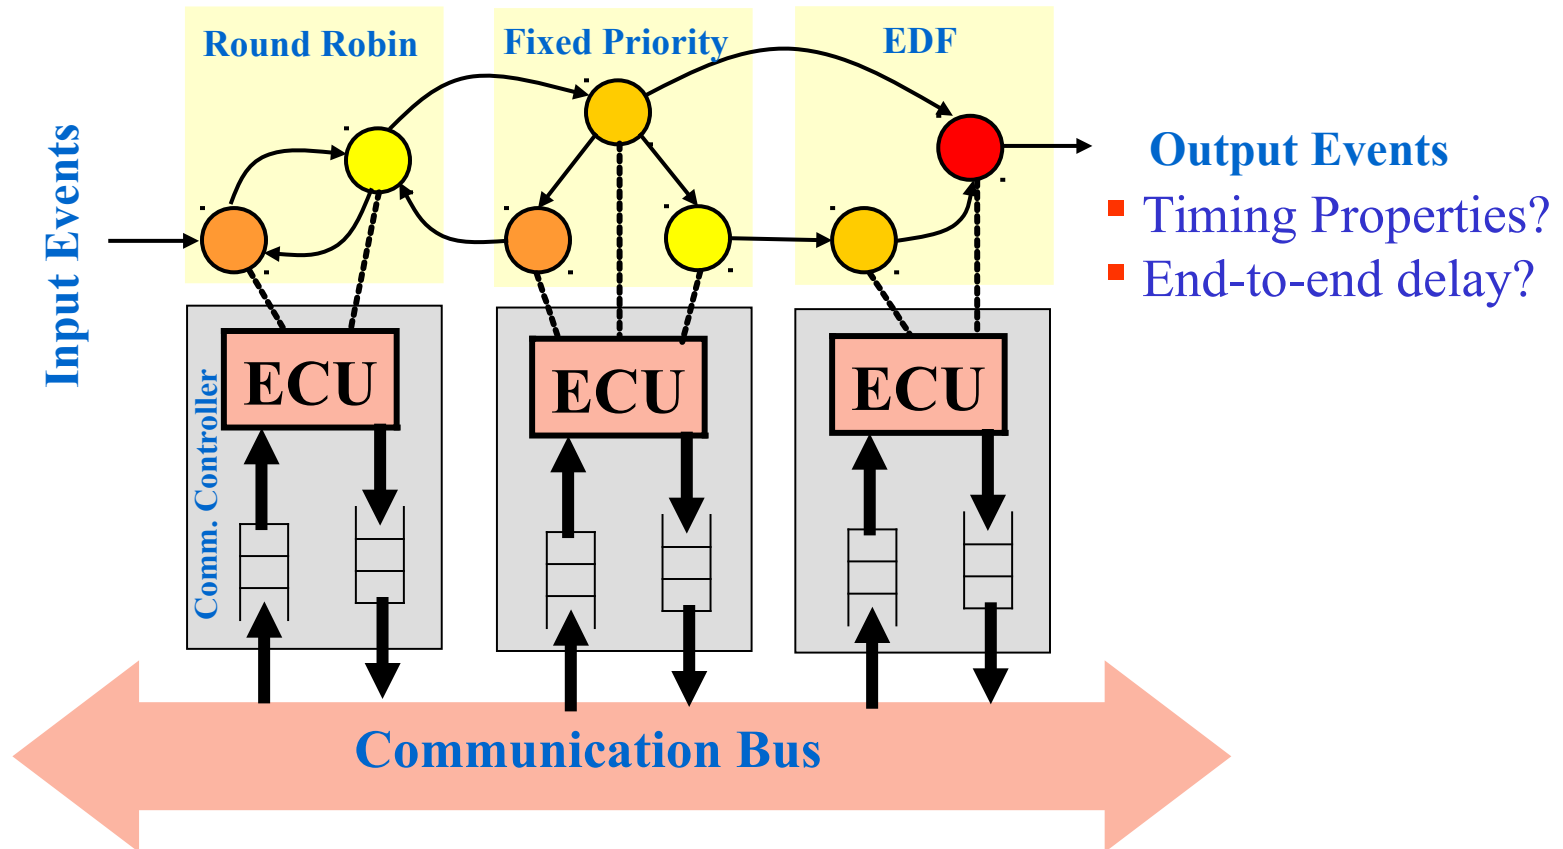
Lecture outline

- Why concurrency?
- Foreground/background vs. multi-tasking systems
- Concurrent processes
 - Communication: message passing vs. shared memory
- Scheduling
- **Bus scheduling**

Bus scheduling

- So far we have studied the scheduling analysis on *one* processor
- However, as systems become more complex, multiple processors exist on a system
- MPSoCs in mobile devices, automotive electronics
- The different processors exchange messages over a *communication bus*!

System-Level Timing Analysis Problem



- Tasks have different activation rates and execution demands
- Each computation/communication element has a different scheduling/arbitration policy

Goal

- Time- and Event-triggered protocols – advantages and disadvantages
 - With the example of bus protocols in automotive systems

Time-Triggered and Event-Triggered Systems



TT



ET

- TT and ET are two students studying at the university
- TT is well-organized and studies according to a predefined schedule
- ET takes life more casually and studies only when necessary

The Story of Two Students – TT and ET




I study Circuit Theory **every Tuesday**

Circuit Theory exam announced!



I have to study Circuit Theory **NOW!**

Circuit Theory exam announced!



The Story of Two Students – TT and ET



I study Circuit
Theory **every**
Tuesday

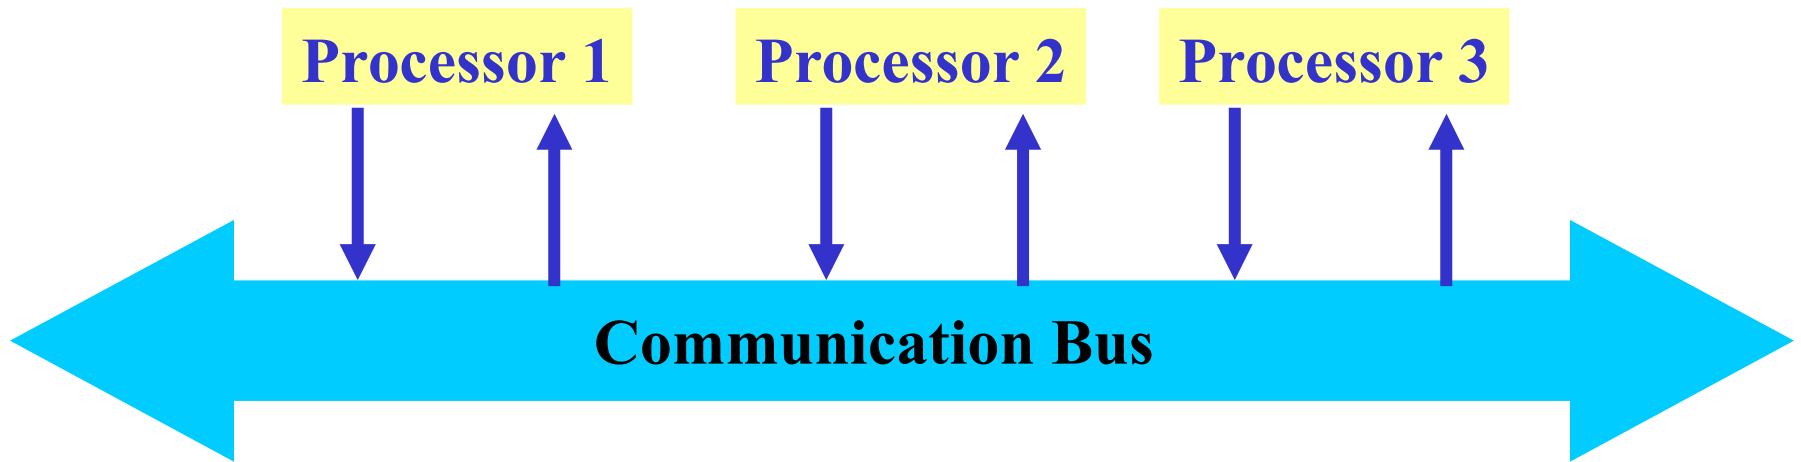
TT follows a
time-driven or Time-Triggered protocol



I have to study
Circuit Theory
NOW!

ET follows a
event-driven or Event-Triggered protocol


Bus Arbitration Policies




- When multiple processors want to transmit data at the same time, how is the contention resolved?
 - Using a bus arbitration policy, i.e. determine who gets priority
 - Examples of arbitration policies
 - **Time Division Multiple Access, Round Robin, Fixed Priority ...**

Time-triggered arbitration


Time-triggered arbitration policy



(Non preemptive) Event-triggered arbitration policy




Bus Arbitration Policies



Time-Triggered Policy

- Only interrupts from the timer are allowed
- Events CANNOT interrupt
- Interaction with environment through polling
- Schedule is computed offline, deterministic behavior at runtime
- Example: Time Division Multiple Access (TDMA) policy

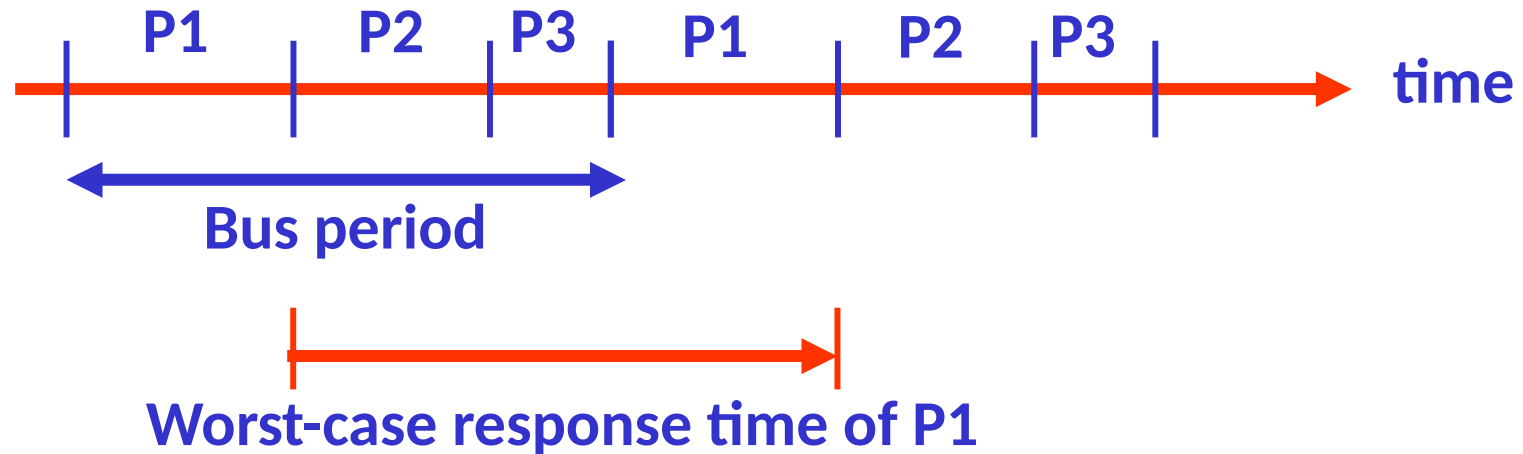
Bus Arbitration Policies



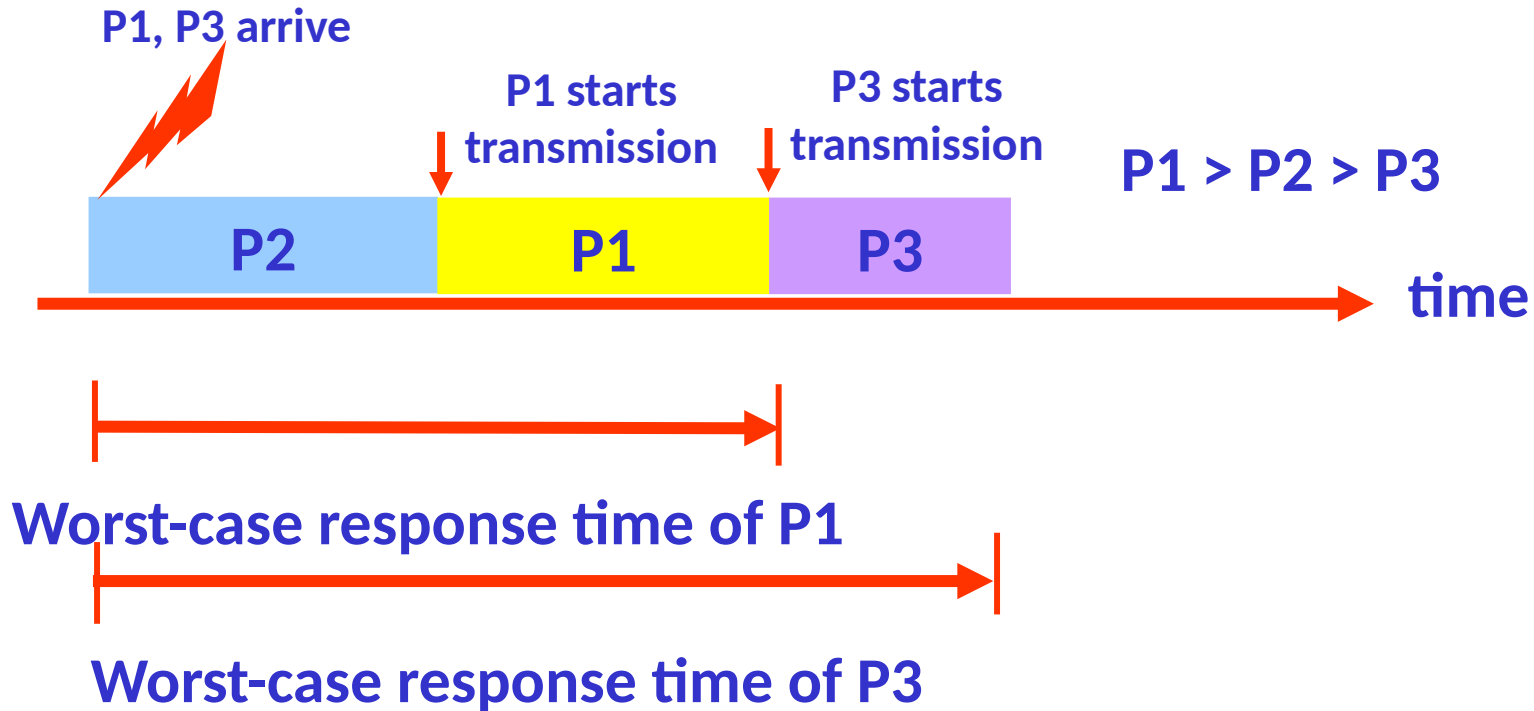
Event-Triggered Policy

- Interrupts can be from the timer or from external events
- Interaction with environment through interrupts
- Schedule is dynamic and adaptive
- Response times 'can' be unpredictable
- Example: Fixed Priority scheduling policy

Computing Response Times in Time-triggered systems



Computing Response Times in Event-triggered Systems



Two well-known bus protocols

- Time-Triggered Bus Protocols
 - **Time-Triggered Protocol (TTP)** – mostly used for reliable/guaranteed communication. Also used in avionics
 - Based on Time Division Multiple Access (TDMA) policy
- Event-Triggered Bus Protocols
 - **Controller Area Network (CAN)** – widely used for chassis control systems and power train communication
 - Based on fixed priority scheduling policy
 - ‘Does not provide’ hard real-time guarantees

Time-Triggered Vs Event-Triggered: Summary

- Both have their advantages and disadvantages
- Recently, combined protocols are being developed

	Time-Triggered	Event-Triggered
Hard Real-Time Guarantees	☺	×
Response Times	×	☺
Bus Utilization	×	☺
Flexibility	×	☺
Composability	☺	×

Conclusion

- TT and ET concepts also applicable to processors
- RMS and EDF are ET or TT?